

The Lessons Register

Every lesson this research has learned, in plain language, and how far each one travels.

Generated 31 August 2026 from the register's own source. Not hand-written, and not hand-editable — the Word file, the Markdown and the code a study calls are three renderings of one thing.

Nothing in this register binds any study. It is a record, not a gate. No standing rule refers to it and no quality gate consults it, deliberately — the walk-forward method that produced its strongest lessons has not itself been validated. The house rule is that nothing enters the method without surviving the same out-of-sample test the forecasts must survive. Read this to think with; do not cite it as authority.

How to read this

A lesson is useless until you know how far it carries. Every entry is therefore tagged with one of three scopes, and the difference is not a matter of taste.

Scope	Means	Who must read it
ALL	About method, arithmetic, or how work gets checked	Every study, every time
CLASS	True of every company that works the same way	The next study of that class only
STOCK	True of this one company and nothing else	That company's next update only

Applying a STOCK lesson to another company is superstition. Applying an ALL lesson is mandatory. The middle category is where most of the value sits, and it is the one that takes longest to fill.

Every lesson also records how it was learned, because the strength of the evidence differs enormously. A walk-forward test measures the method against real outcomes across many starting points and is the strongest evidence this research produces; an outside critique, a self-audit and a defect found while building are each weaker in turn. Where two lessons disagree, the walk-forward one wins.

And every lesson carries what would overturn it. A lesson with no falsifier is a habit, not a finding, and habits are how a method quietly stops being tested.

What is in here, and what is honestly missing

	Lessons
Bind on every study (ALL)	33
Bind on a class of company (CLASS)	14
Bind on a single company (STOCK)	5
Total	52

From fundamental walk-forward testing	14
From price-engine walk-forward testing	2
From outside critiques	10
From self-audits	7
Found while building	19

Two different tests are both called a walk-forward

They test different machinery on different evidence, and the first edition of this register conflated them — which understated the evidence base badly and is recorded here rather than quietly corrected.

	What it tests	Names	Resolved forecasts
Fundamental	The forecasting method — project each driver from a past origin, score revenue, cost and profit against what happened	1 (PHDC)	10 origins x 5 horizons
Price engine	The probability cone — strike it at a past origin and score band coverage and a proper score against a naive rule	19	317

The price engine is well tested; the fundamental method is not. 19 names carry price-engine evidence, DU (18 forecasts) and GBCO (17) among them. The fundamental method has been through a full training run on PHDC alone, and that run's own record states its corrections rest on two starting points, its intervals are wide with several straddling zero, and its observations are not independent. Every lesson from the fundamental method is therefore marked PROVISIONAL; price-engine lessons are not, because they rest on 317 forecasts across 19 names.

Not yet acted on (16): L-001 — Revenue and cost must be recognised on the same clock; L-002 — A correction factor is only honest when the model is right; L-003 — Before adopting any correction, ask whether it matches how we treat that driver everywhere else; L-010 — Never divide one year's value by another year's volume; L-011 — Publish long-horizon forecasts as ranges, never as points; L-012 — Company guidance is scored, never used as an input; L-101 — Volume is set by the launch calendar, and no demographic anchor can see it; L-102 — Check the revenue-recognition basis before anything else; L-104 — Deliveries must be constrained by what has actually been sold; L-203 — Palm Hills' 2025 balance sheet and cash-flow statement disagree by 47% of revenue; L-204 — Palm Hills' interest rate must be measured on the borrowings that actually bear interest; L-205 — Palm Hills' operating drivers stop a year before its financial statements do; L-028 — Depreciation forecast from an asset base that is itself forecast compounds its own error; L-114 — A developer's new-sales value is under-forecast whenever price and volume are projected separately; L-029 — A bias that changes direction between regimes must never be corrected for; L-030 — Test every measured bias for stability across regimes before acting on it. Recorded as open rather than quietly carried as done.

Lessons that bind on EVERY study

Read these before starting any study, of any company, in any market.

L-001 Revenue and cost must be recognised on the same clock.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

If you count the money coming in as the building goes up, you must count the money going out the same way. Book revenue as you build but costs only when you hand over the keys, and the costs arrive late — so every year looks more profitable than it is.

How we know. Gross profit came out about 72% too high, in 86% of tested cases, and that fed through into a net profit about 3x too high in 97% of cases — worse than simply assuming last year's profit repeats.

What would overturn it. An issuer that genuinely recognises revenue only on handover, where the two clocks already coincide.

L-002 A correction factor is only honest when the model is right.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

When a forecast is consistently wrong in the same direction, it is tempting to add a fudge factor. But a steady error usually means something is wired wrong, and a fudge factor hides the wiring instead of fixing it. Fix the wiring.

How we know. Interest was being divided by total current liabilities (EGP 105.1bn) instead of by debt that actually pays interest (EGP 24.1bn) — a denominator 4.4x too big, implying a 3.19% borrowing rate for a company that borrows at 13.91%. The resulting 'bias' looked real and a half-strength correction 'worked', cutting the error from 0.85 to 0.40.

What would overturn it. Nothing. A mis-specification is never fixed by a multiplier.

L-003 Before adopting any correction, ask whether it matches how we treat that driver everywhere else.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

A number that is out of line with the rest of the book usually means our own method slipped on this one name — not that the company is unusual. That second question is what tells the two apart.

How we know. The finance-cost correction passed its own test and was blocked by this one, which is how the wrong denominator in L-002 was found. Every other Egyptian study builds interest from named facilities, each at its own rate.

What would overturn it. Nothing yet. If a correction ever fails this test for a reason that turns out to be a genuine company difference, record it here.

L-004 Country risk must be counted once, not twice.

EVERY STUDY applies to every study · learned from a self-audit, GBCO study; corrected across the book

A local government bond yield already contains the country's own risk. Add a country-risk premium on top of it and you have charged for the same danger twice, which makes every company look worth far less than it is.

How we know. The error reached a delivered study and forced a re-issue of every affected valuation. The fix: normalise the risk-free rate by subtracting that sovereign's own default spread before adding the premium back.

What would overturn it. A market where the quoted government yield genuinely carries no default component.

L-005 Margins are an output, never an input.

EVERY STUDY applies to every study · learned from a outside critique, DU study, 17-Aug-2026

If the filings tell you what things cost per unit, build the cost from those units and let the margin fall out. Typing in a margin means you have assumed the answer to the most important question in the model.

How we know. Edition 3 held one contribution margin per segment as an input at the audited rate, discarding six months of reviewed data. Rebuilding it bottom-up surfaced a downside nobody had found: a flat blended ARPU was a mix tailwind masking per-leg erosion, worth -17.3% on the cash-flow lens.

What would overturn it. An issuer whose filings genuinely disclose no unit or per-segment cost detail — in which case say so, and flag the gap.

L-006 A self-audit that only re-checks what you did will keep missing what you never did.

EVERY STUDY applies to every study · learned from a outside critique, DU study, 17-Aug-2026

Checking your own sums proves the sums are right. It does not tell you about the whole disclosure you never opened. Ask what the filings contain that the model does not use.

How we know. DU's self-audit found seven defects no critique had raised, and still missed the largest one — that the cost side had not been built from the bottom up at all. It surfaced only when a rebuild was forced.

What would overturn it. Nothing. This is a rule about how to audit, not about any company.

L-007 Try the company's own website before any data provider.

EVERY STUDY applies to every study · learned from a outside critique, ARCC study, 07-Aug-2026

The company's own investor-relations page is the source. A data provider's summary is somebody else's transcription of it, and transcriptions drift.

How we know. Adopted after a study leaned on aggregator figures while the company's own quarterly filing sat unread. Log every attempt, including the ones that fail — a site that will not load is a fact worth recording.

What would overturn it. Nothing. Where the primary source is unreachable, stop and say so rather than substituting a weaker one.

L-008 A period is not researched until both its statements AND its results release are in.

EVERY STUDY applies to every study · learned from a outside critique, MODON study, 09-Aug-2026, found by an outside critique

The financial statements give you the numbers. The results release gives you what they mean — the order book, the volumes, the sales. Reading one without the other leaves the model running on stale drivers.

How we know. MODON's first edition read the half-year statements but not the release published the same month, whose AED 65.4bn order book and AED 26bn half-year sales superseded the study's core drivers. Caught externally, restruck the same day.

What would overturn it. An issuer that publishes no results release at all — record the absence.

L-009 One escalator per cost driver, never one blended rate across all of them.

EVERY STUDY applies to every study · learned from a outside critique, ARCC study, 07-Aug-2026

Coal, imported fuel, local wages and transport do not all rise at the same speed. Escalating them together with a single inflation number invents a margin trend that is an artefact of the shortcut.

How we know. The whole of one study's forecast margin decline turned out to be a mechanical consequence of setting the price path below a single blended cost index every year — and was already contradicted by a realised quarterly margin above the prior full year.

What would overturn it. A cost base genuinely made of one input.

L-010 Never divide one year's value by another year's volume.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

If sales are disclosed to 2024 and unit counts only to 2023, dividing one by the other does not give you a price. It gives you a number with no meaning.

How we know. Taking each series' latest available year separately produced a price per unit of 28.5 against a true figure of 11.2 — off by a factor of two and a half.

What would overturn it. Nothing. This is arithmetic.

L-011 Publish long-horizon forecasts as ranges, never as points.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

A five-year forecast is not one number. Show the spread the method has actually produced in testing, so a reader can see how much is known.

How we know. At five years the measured spread on revenue ran from 83,620 to 214,090 EGP mn, and rested on five resolved observations. Publishing a single figure would have implied precision that ten origins cannot support.

What would overturn it. A record long enough and tight enough that the range collapses — which would itself be the evidence.

L-012 Company guidance is scored, never used as an input.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

Management's forward targets lean the same way an optimistic model does. Feeding them in imports the bias instead of correcting for it. Grade them against outcomes instead.

How we know. Every target quoted retrospectively had been beaten; both targets that could be graded BEFORE the outcome were missed, over-forecasting handovers by about 25%. The model's own delivery bias leaned the same way.

What would overturn it. An issuer with a long record of forward targets graded in its own disclosures.

L-013 A recent reviewed actual beats a stale full-year rate.

EVERY STUDY applies to every study · learned from a outside critique, DU study, 17-Aug-2026

Anchor unit costs and rates on the most recent reviewed period. Only let them drift where a named mechanism has a measured direction in the company's own figures — and hold everything else flat, including improvements you would like to believe.

How we know. Carrying a first-half rate into the second half was proved to OVERSTATE cost, using the prior year's own actual halves.

What would overturn it. A structural mechanism with a measured like-for-like direction in the issuer's own period pair.

L-014 A check must look at the thing that actually ships.

EVERY STUDY applies to every study · learned from a found while building, PHDC study build, 30-Aug-2026

It is easy to build a check that inspects a slightly different object from the one the reader sees — and it will report clean forever.

How we know. The table-width audit read cell widths back and reported clean, while the renderer reads a separate column grid that nothing had set. MEASURED ACROSS THE BOOK RATHER THAN ASSUMED: 39 of 984 tables in 44 delivered documents have a grid disagreeing with their stated widths, concentrated in two (31 of 33, and 8 of 10). Most builders escape it only because they happen to set COLUMN widths, which the library propagates into the grid — an accident of which call was used, not something any check was testing. The same species had already appeared twice: a beta check that trusted a self-set flag, and page checks that parsed with a regex where the browser reads the last value.

What would overturn it. Nothing. When this species appears again, close the class, not the instance.

L-015 An empty result is not a clean result.

EVERY STUDY applies to every study · learned from a found while building, Repo-wide, 25-Aug-2026; recurred in the PHDC build

A check that examined nothing looks exactly like a check that found nothing wrong. Always count what was examined against a total held somewhere else.

How we know. Four instances in one session, then a fifth: the study's four figures had no committed code that built them at all. They were produced once by hand and thereafter only checked, so one figure was drawing a grid the model had stopped producing and another could not be rebuilt at all.

What would overturn it. Nothing. When a probe comes back empty, first assume the probe did not run.

L-016 One document, one model.

EVERY STUDY applies to every study · learned from a found while building, PHDC study build, 30-Aug-2026

If two different models feed one report, it will publish two different answers and nobody will notice, because each one looks right beside its own neighbours.

How we know. The headline range, the sensitivity grid and the entire workbook came from a ten-year capacity model while the summary tables and the bridge came from the bottom-up model — EGP 9.17-38.74 in the headline against EGP 2.85-38.75 two pages later. Found by independently recalculating the delivered workbook, not by reading it.

What would overturn it. Nothing. A deliberately published second model must be labelled as a cross-check, never as the answer.

L-017 Never resolve a spreadsheet row by its position.

EVERY STUDY applies to every study · learned from a found while building, PHDC study build, 30-Aug-2026

Row numbers move the moment anyone inserts a line. A check pointed at the wrong row still computes, and still reports a number. Look rows up by their label.

How we know. After one driver was added to the workbook, the recalculation read 'revenue per delivered unit' where it expected revenue and reported a mismatch four orders of magnitude wide.

What would overturn it. Nothing.

L-018 A registered input that no line of the model consumes is a defect.

EVERY STUDY applies to every study · learned from a self-audit, AMOC study self-audit, 06-Aug-2026

Putting a number in the register is not the same as using it. A cost the company really pays, sitting in the register and consumed by nothing, is money the valuation has quietly ignored.

How we know. Employees' profit share and board bonuses — about EGP 135mn a year, a real cash distribution — were registered, quoted in the delivered documents, and consumed by no line of the model. Worth about -7.7% of the central value.

What would overturn it. An input registered deliberately as context, explicitly marked as not consumed.

L-019 A balance sheet that closes on a residual is not a balance sheet.

EVERY STUDY applies to every study · learned from a self-audit, SCEM study self-audit, 06-Aug-2026

If total liabilities are computed as whatever makes the sheet balance, then every mistake anywhere else disappears into that one line and the sheet balances no matter how wrong it is.

How we know. One study wrote total liabilities as assets minus equity. The residual fell to about 706 against a disclosed prior-year figure of 1,611, and nothing in the model objected.

What would overturn it. Nothing. Where a residual is genuinely unavoidable it must be a named line with a stated meaning, not a silent plug.

L-020 Bear and bull cases must come from the model, not from multipliers.

EVERY STUDY applies to every study · learned from a self-audit, SCEM study self-audit, 06-Aug-2026

A 'bear case' produced by multiplying the central figure by 0.88 is not a scenario. It is decoration, and presenting it in a table of weighted outcomes misrepresents how it was made.

How we know. The published bull case was literally the central value times 1.14 and the bear times 0.88, presented as a weighted range that no weighting could produce. A separate study's published range turned out to be the minimum and maximum of one lens's extremes.

What would overturn it. Nothing.

L-021 Never fit an input to get the output you wanted.

EVERY STUDY applies to every study · learned from a self-audit, SCEM study self-audit, 06-Aug-2026

Introducing a number specifically to stop a line going negative is not an estimate. It is the answer working backwards into the assumptions.

How we know. A house estimate was introduced specifically to keep one year's EBITDA from going negative under an earlier scaling, and then carried as though it were sourced.

What would overturn it. Nothing.

L-022 A recognised provision is a probable outflow and belongs in the bridge.

EVERY STUDY applies to every study · learned from a self-audit, AMOC study self-audit, 06-Aug-2026

If the company has formally recognised that it probably owes something, the valuation cannot leave it out on the grounds that it is disputed.

How we know. A recognised tax-disputes provision of EGP 904.6mn was excluded on the reasoning that the discounted earnings were struck after tax — which confuses a forward tax rate on current profits with settling past assessments. Worth -9.8% of the central value.

What would overturn it. A provision the issuer has explicitly reversed or settled.

L-023 A beta must be regressed against the exchange's published index.

EVERY STUDY applies to every study · learned from a found while building, Repo-wide; measured on FERTIGLB, 10-Aug-2026

Comparing a share against a basket of the other shares you happen to cover is not comparing it against the market. The basket changes every time you add a company.

How we know. Every study in the repository had regressed against an equal-weighted composite of covered names, because the first one did and each copied the last. On the first name run both ways the composite gave a beta of 0.492 against the real index's 0.931 — a 40% understatement that carried the cost of capital from 11.90% to 8.53% and overstated value by 21.6%.

What would overturn it. Nothing. Where the exchange's index is not held, stop and ask for it rather than building a substitute.

L-024 A cautious-sounding label is still a claim, and gets audited like one.

EVERY STUDY applies to every study · learned from a found while building, Repo-wide, 24-Aug-2026

Understating in the wrong direction is not a safe error. Conservative language slips through review precisely because it does not sound like a boast.

How we know. Pages carried the label 'failed calibration test' on names whose forecast bands had actually been too WIDE — over-covering at 97% against a 90% target — while every name whose bands genuinely ran narrow carried no flag at all. The label pointed at the wrong names for weeks.

What would overturn it. Nothing.

L-025 A test that has never changed an outcome is decorative.

EVERY STUDY applies to every study · learned from a found while building, Repo-wide, 25-Aug-2026

Before keeping a check, ask what it has ever rejected. If the answer is nothing, it is measuring the wrong thing or sitting at the wrong threshold — and it is occupying the place a working check should hold.

How we know. One standing gate had never excluded anything: across every verdict ever recorded the tally was 213 / 94 / 12, with no market ever blocked, while both governing documents described it as the standing gate. It was retired.

What would overturn it. Nothing.

L-026 Two figures for the same thing must both be published, never averaged.

EVERY STUDY applies to every study · learned from a found while building, Repo-wide; applied to PHDC's cash-flow wedge, 30-Aug-2026

Where a number has two legitimate readings, showing the midpoint hides that there was a disagreement at all. Show both and say which one the answer rests on.

How we know. PHDC's audited balance sheet and audited cash-flow statement disagree about one year by 46.8% of revenue, and the gap cannot be split from what is disclosed. The forecast is published both ways rather than reconciled by an invented plug.

What would overturn it. Disclosure that resolves the disagreement.

L-027 Never blend two legs that need different methods.

EVERY STUDY applies to every study · learned from a found while building, House method, standing protocol

A company with a manufacturing arm and a lending arm is two businesses. Forcing them into one model produces a number that describes neither.

How we know. Each class has a method that fits it: a developer on sum-of-the-parts, a bank on dividends and returns on equity, a holdco piece by piece, a contractor on cash flow plus parts. Where a company spans two, the legs are valued separately and added.

What would overturn it. Legs so similar that one method genuinely fits both — stated, not assumed.

L-028 Depreciation forecast from an asset base that is itself forecast compounds its own error.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

Depreciation is usually projected as a rate on a fixed-asset balance that the model has also projected. Two forecasts stacked on each other drift much further than either one alone.

How we know. Bias +0.586 log (about 1.8 times too high), average miss 1.112, wrong in the same direction in 77% of cases, and the sign holds across every bootstrap block tested (n=39).

What would overturn it. A build where depreciation is taken from a disclosed schedule rather than from a projected asset base, and the error stays large.

L-029 A bias that changes direction between regimes must never be corrected for.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

If a driver runs high in one period and low in the next, the average is a number that was never true. Record the instability and leave the driver alone.

How we know. By era: E1 pre-float -0.315; E2 post-float +0.089; E3 devaluation -0.015.

What would overturn it. A record long enough that one sign dominates across every regime tested.

L-030 Test every measured bias for stability across regimes before acting on it.

EVERY STUDY applies to every study · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

A bias measured over a whole history can hide two opposite halves. Split the record by regime and check the sign holds in both before you treat it as a fact.

How we know. By era: E2 post-float +0.507; E3 devaluation -0.196.

What would overturn it. Nothing. This is a rule about how to read a measured bias.

L-031 Band coverage and a skill score answer different questions, and can disagree.

EVERY STUDY applies to every study · learned from a price-engine walk-forward test, Price-engine walk-forward, 19 names

A forecast range can hold exactly as often as promised while still being no better than a simple rule of thumb. Both facts are true at once. Ask which question you are answering before you quote either number.

How we know. Across 317 resolved three-month forecasts on 19 companies, the 90% band caught the outcome 94% of the time and the 80% band 86% — close to what was promised — while the same forecasts beat a carry-anchored random walk on a proper score only 44% of the time. Coverage said the cone was honest; the score said it was no better than the naive rule.

What would overturn it. A book where the two measures agree closely across most names, which would mean one of them had stopped being informative.

L-032 A book-level average hides how differently individual names behave.

EVERY STUDY applies to every study · learned from a price-engine walk-forward test, Price-engine walk-forward, 19 names

Averaging a measure across every company can look reassuring while individual names are far off in both directions. Report the spread, not just the average.

How we know. The middle band caught the outcome half the time on average, as designed — but per name it ran from 20% to 92% against the same 50% target. The average was right and told you almost nothing about any particular company.

What would overturn it. A book whose per-name spread narrows enough that the average is representative — which would itself have to be measured, not assumed.

L-033 Check the column name before believing a column of zeros.

EVERY STUDY applies to every study · learned from a found while building, Price-engine walk-forward, 19 names

A measurement that comes back the same for everything is more often a broken query than a real finding. Look at the actual field names before you conclude anything from a suspiciously uniform answer.

How we know. Scoring the whole book against its benchmark returned 0% for 18 of 19 names and 53% for the nineteenth. The 18 files name the column crps_b and the one name crps_bench; the query asked for crps_bench. The real answer is 44%. This is the same species as L-015 and it recurred inside the very session that recorded it.

What would overturn it. Nothing. When a probe comes back uniform or empty, the first hypothesis is that the probe did not run.

Lessons that bind on a CLASS of company

Read the section for the class being studied. Do not read across classes — a rule true of developers is not evidence about airlines.

Real-estate developer, off-plan, percentage-of-completion

L-101 Volume is set by the launch calendar, and no demographic anchor can see it.

A CLASS OF COMPANY applies to every real-estate developer, off-plan, percentage-of-completion · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

How many units a developer sells next year depends on which projects it chooses to launch. Population growth cannot predict that. Use a population anchor to keep the driver honest, but say out loud that it runs low.

How we know. Units sold were under-forecast by about 19%, consistently, and the miss widened sharply during the devaluation years. One region's annual unit count ran 726, 1,552, 800, 1,231, 4,192 — not a trend, a launch schedule.

What would overturn it. A developer that discloses a firm forward launch pipeline, where volume can be built from the pipeline instead.

L-102 Check the revenue-recognition basis before anything else.

A CLASS OF COMPANY applies to every real-estate developer, off-plan, percentage-of-completion · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

Developers differ in when they book a sale. Get this wrong and every profit line in the model is wrong, no matter how carefully the rest is built.

How we know. This is the class-specific form of L-001. Palm Hills moved to percentage-of-completion in January 2016; a model built without noticing produced a net-profit forecast three times too high.

What would overturn it. Nothing for this class — but confirm the basis per issuer rather than assuming it.

L-103 A developer holds two contract positions at once, and they must move together.

A CLASS OF COMPANY applies to every real-estate developer, off-plan, percentage-of-completion · learned from a found while building, PHDC study, 30-Aug-2026

Buyers who have paid ahead sit as a liability; buyers still paying instalments sit as an asset. Both grow out of the same order book, so modelling one against revenue and the other against the order book makes them drift apart and the cash forecast falls apart.

How we know. Driving receivables off revenue and customer advances off the order book produced a projection with cash at minus EGP 171bn and a negative present value for the forecast period. Both legs are contract-linked and must be driven consistently.

What would overturn it. A developer that sells for cash on completion, where only one leg exists.

L-104 Deliveries must be constrained by what has actually been sold.

A CLASS OF COMPANY applies to every real-estate developer, off-plan, percentage-of-completion · learned from a found while building, PHDC study, 30-Aug-2026 · **STATUS: OUTSTANDING**

A developer cannot hand over more homes than it has contracted. A delivery path that compounds on its own can quietly outrun the order book.

How we know. The delivery driver compounds at 15% a year with nothing linking it to the order book. It happens not to bind on the current path, but it is unguarded.

What would overturn it. Nothing.

L-105 For a developer, the collection cycle and a fast growth path cannot both hold.

A CLASS OF COMPANY applies to every real-estate developer, off-plan, percentage-of-completion · learned from a found while building, PHDC study, 30-Aug-2026

If buyers pay over three years and revenue is growing 40% a year, the money going out runs far ahead of the money coming in. Something has to give, and the study should say which.

How we know. Holding the collection cycle produces free cash flow negative in all five years and needs EGP 168.5bn of new borrowing by 2030, whose interest alone exceeds operating profit. Holding cash conversion instead requires the collection period to fall from 948 days to 299.

What would overturn it. A developer collecting substantially at or before handover.

L-114 A developer's new-sales value is under-forecast whenever price and volume are projected separately.

A CLASS OF COMPANY applies to every real-estate developer, off-plan, percentage-of-completion · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

Sales value is units times price. Project each one conservatively and the two shortfalls multiply, so the value comes out much lower than it should. Check the product, not just the parts.

How we know. Bias -0.369 log (about 45% too low), average miss 0.480, wrong in the same direction in 66% of cases, and the sign holds across every bootstrap block tested (n=35).

What would overturn it. A developer where price and volume are projected jointly and the value forecast is still biased low.

Telecom operator

L-106 A flat blended revenue-per-user can hide two opposite movements.

A CLASS OF COMPANY applies to every telecom operator · learned from a outside critique, DU study, 17-Aug-2026

An average that does not move can be a growing share of expensive customers cancelling a falling price on each one. It looks stable right up until the mix stops shifting.

How we know. A flat blended ARPU turned out to be a +2.6% mix tailwind against -2.4% per-leg erosion — worth -17.3% on the cash-flow lens once the tailwind exhausts, and invisible until the average was taken apart.

What would overturn it. An operator with a stable customer mix, demonstrated rather than assumed.

L-107 Where a company discloses one total two different ways, recover the cross-tabulation and test it.

A CLASS OF COMPANY applies to every telecom operator · learned from a outside critique, DU study, 17-Aug-2026

If costs are disclosed both by type and by division, those two views together often pin down the split the company never publishes — but the recovered figures must add up exactly to the disclosed lines in every period, or the assumptions behind them are wrong.

How we know. Direct costs disclosed three ways by nature and four by segment let the joint be recovered exactly. The test is that every residual is economically possible and foots to the disclosed line in every disclosed period, asserted in code.

What would overturn it. A recovered residual that comes out negative or fails to foot — in which case the assumption set is wrong and the joint must not be used.

Cement and heavy industrial

L-110 Fuel and imported inputs escalate on their own commodity path, not on domestic inflation.

A CLASS OF COMPANY applies to every cement and heavy industrial · learned from a outside critique, ARCC study, 07-Aug-2026

A globally traded input like coal follows the world price and the exchange rate. Escalating it with the local cost of living invents a margin story.

How we know. The class-specific form of L-009, and the case it was adopted from: a reconciliation against an outside broker's model showed the entire forecast margin decline was an artefact of one blended index.

What would overturn it. An input genuinely priced domestically.

Petrochemical

L-109 A forecast year must be producible from the disclosed half-year.

A CLASS OF COMPANY applies to every petrochemical · learned from a outside critique, BOROUGE study, 17-Aug-2026

If hitting your full-year number requires a second half better than any full year the company has ever had, the forecast is not a forecast.

How we know. The 2026 forecast needed a second-half margin of 43.2% against a best-ever full year of 41.0%, and a feedstock cost below the model's own floor. Two independent critiques reached this separately.

What would overturn it. Disclosed evidence of a genuine step-change in the second half.

Airline

L-108 Leased aircraft must carry a liability, a rental and a cash cost.

A CLASS OF COMPANY applies to every airline · learned from a self-audit, AIRARABIA study self-audit, 10-Aug-2026

Adding leased planes to the asset side without the obligation that comes with them makes the airline look like it got them free.

How we know. AED 300mn a year of leased-fleet additions entered invested capital with no incremental lease liability in the bridge, no rental in the cost stack and no cash outflow in free cash flow. Direction: overstates.

What would overturn it. An issuer that owns its fleet outright.

Bank

L-111 A bank is valued on what reaches the shareholder, not on enterprise value.

A CLASS OF COMPANY applies to every bank · learned from a found while building, House method, ADCB pattern

For a bank, debt is raw material rather than financing, so the usual enterprise-value approach is meaningless. Value it on dividends, on free cash flow to equity, and on returns above the cost of equity.

How we know. The reference bank study runs a dividend-discount model alongside free cash flow to equity and residual income. An enterprise-value discounted cash flow on a bank subtracts its own deposits as though they were borrowings.

What would overturn it. Nothing for a deposit-taking institution.

Holding company

L-112 A holding company is valued piece by piece, never as one consolidated business.

A CLASS OF COMPANY applies to every holding company · learned from a found while building, House method, ALPHADHABI pattern

A holdco is a collection of stakes in different businesses, each worth what it is worth on its own terms. Merging them into one cash-flow forecast averages away exactly the information you need.

How we know. The reference holdco study values each holding separately and sums them, with the discount to that sum stated rather than assumed. Ownership stakes are sourced to the named transaction, never estimated.

What would overturn it. A holdco whose subsidiaries are so alike that one method genuinely fits them all — in which case say why.

Commodity and metals

L-113 Say plainly when a commodity's forecast rests on borrowed calibration.

A CLASS OF COMPANY applies to every commodity and metals · learned from a found while building, House method, standing protocol

Some metals have no history of their own in this system and borrow another metal's settings. That is a real limitation, and a reader is owed it in plain words rather than a footnote.

How we know. Gold is calibrated on its own data alone, so testing it against that same data proves nothing. Silver is published with no fit of its own at all — it borrows gold's. Neither may be presented with the confidence of a name that has a proper panel behind it.

What would overturn it. Enough silver, copper and platinum history to fit each on its own.

Lessons that bind on ONE company

Read only the section for the company being updated. These do not generalise and must not be applied to another name.

PHDC

L-201 Palm Hills' units-sold figure does not affect its valuation.

ONE COMPANY ONLY applies to PHDC only · learned from a found while building, PHDC study, 30-Aug-2026

Revenue comes from homes handed over, not homes sold. New sales only move the reported order book. Worth knowing before spending time refining the sales forecast.

How we know. Raising units sold by 50% moves value per share by exactly zero. The driver that does move it is the delivery growth rate.

What would overturn it. A model where the order book constrains deliveries — which is exactly what L-104 asks for.

L-202 Palm Hills' regional unit charts overlap and must be reconciled before use.

ONE COMPANY ONLY applies to PHDC only · learned from a found while building, PHDC study and walk-forward, 30-Aug-2026

Adding up the regional figures does not necessarily give the company total, because the charts double-count. Check each release against its own printed total before trusting the sum.

How we know. Only releases whose regional blocks reconcile to their own group total within 2% are accepted. Where they do reconcile, the regional sum runs 0.8% to 1.5% above the printed total — rounding, not a different measure.

What would overturn it. Nothing; re-check on each new release.

L-203 Palm Hills' 2025 balance sheet and cash-flow statement disagree by 47% of revenue.

ONE COMPANY ONLY applies to PHDC only · learned from a found while building, PHDC study, 30-Aug-2026 · **STATUS: OUTSTANDING**

The two audited statements tell different stories about how much cash the year absorbed, and the published cash-flow statement is too summary to explain the difference.

How we know. The balance sheet says working capital rose EGP 20,084.7mn; the cash flow implies EGP 3,146.2mn. The gap of EGP 16,938.5mn cannot be split from three published totals, so the forecast is published both ways.

What would overturn it. The full 2025 cash-flow statement, line by line.

L-204 Palm Hills' interest rate must be measured on the borrowings that actually bear interest.

ONE COMPANY ONLY applies to PHDC only · learned from a fundamental walk-forward test, PHDC walk-forward and study, 30-Aug-2026 · **STATUS: PROVISIONAL**

Most of what the company owes — customer deposits, suppliers, cheques — pays no interest at all. Measuring the borrowing rate against everything it owes makes it look four times cheaper than it is.

How we know. Total current liabilities of EGP 105,099mn against EGP 24,069mn that bears interest. The correct rate is 13.91%, not 3.19%. Part of the charge is also capitalised into work in progress rather than expensed.

What would overturn it. A change in the disclosed debt structure.

L-205 Palm Hills' operating drivers stop a year before its financial statements do.

ONE COMPANY ONLY applies to PHDC only · learned from a fundamental walk-forward test, PHDC walk-forward, 30-Aug-2026 · **STATUS: PROVISIONAL**

The company files accounts for a year before it publishes the results release that carries units and sales — so the volume anchor is always one to two years behind the revenue figure.

How we know. Units sold stop at FY2023 and new sales at FY2024 while revenue runs to FY2025. Every affected cell records the lag rather than filling it.

What would overturn it. The company publishing a full-year results release for the latest filed year.

How a lesson gets added

A walk-forward run's own results are read mechanically and the candidate findings are drafted with their evidence already filled in from what the run measured — so a lesson's 'how we know' clause cannot drift from the numbers behind it.

Choosing the scope is not automated, deliberately. It is a judgement, and getting it wrong is costly in both directions: too narrow and the next study repeats the mistake, too broad and one company's quirk becomes a house rule nobody can dislodge. Every draft therefore arrives unscoped and nothing enters this register until someone decides. When genuinely unsure, file at the narrower scope and widen when a second company shows the same thing — one observation is not a pattern.

Every drafted finding ends either registered or declined with a reason. A finding nobody ruled on fails the check, because an unanswered question must not be able to pass as a clean result.