

Fundamental recalibration — the process, end to end

ONE NAME AT A TIME, never in parallel — a defect found on the first usually changes how the second is built. Publishing is a separate gate and it can refuse.

0	FREEZE, THEN ASK	Freeze fair{ } BEFORE anything moves it — a baseline taken afterwards is a fabricated zero. Ask for the latest price and its date, then route around it. Read what already binds on this name and its class.	
1	RECALIBRATE THE FUNDAMENTALS	Walk-forward on the company's own history. Drivers ground up — volume x price, cost per unit, margins as OUTPUTS. Terminal on a DISCLOSED asset life. One class primary IS the central; there is no blend.	
2	INTERNAL QC	The QC gate filled from OUTSIDE the study, evidence per row, never self-certified. Every CI gate run as CI runs it. More than 10% from price either way — the eight-heading gap review, BEFORE anything is staged.	
3	EXTERNAL NEWS RESEARCH	The session BUILDS the prompt — registered name in English and the local language, the study's own driver headings, its dated negative searches — and STOPS. What comes back is a LEAD and never an input; every claim is traced to the primary source before it moves anything.	STOPS — HANDS OFF TO YOU
4	DOES IT MOVE A DRIVER	NO: say so, record which rings were re-run, move on — the commonest outcome. YES: score the prior study first, rebuild only what moved, and move every record, field name and sentence the lever touched IN THE SAME PASS, PDFs re-rendered.	
5	EXTERNAL AUDIT	The session hands you the document and the workbook and STOPS. Findings come back priced before judged, one row each, rejected only with receipts. Where no outside audit is run, the QC auditor stands in — from outside the study.	STOPS — HANDS OFF TO YOU
6	THE PUBLISH GATE	Three tests, all must pass. THE GAP: held only more than 10% BELOW price, released by a filed dissent. THE METHOD HOLD: Phase 1 must be proven. READABILITY: an unreadable study is held too. Publishing is a separate, explicit ask.	

The two red steps STOP and wait for you — they are where a 90-name programme actually queues. Expect step 6 to say HELD: most of the book is held on the METHOD until Phase 1 closes, and that is the gate working.

Ninety names through it — who does what

The page opposite is ONE name. This is the part that was missing: which name is touched next, what it is waiting for, and which of us is holding it.

WHAT YOU DO — THREE THINGS

ONE. Say "start a recalibration on {TICKER}" — or "what is next" and take the name the board gives.

TWO. Answer the price question at the START. One line: the latest price and its date. It is the one input only you can supply; the work routes around it rather than waiting.

THREE. Take the two hand-off packs — the external news prompt at step 3 and the document plus workbook at step 5 — and bring back the answers.

Nothing else is yours. You never type a command.

WHAT THE SESSION DOES

Reads the board, takes the next name in wave order, works steps 0 to 6 IN ORDER, stops at each stop condition and reports before moving on.

Stops dead at 3 and 5 and hands over. Never runs two names at once.

Never publishes without an explicit ask, never moves a fair value toward a price, never edits a gate to make something pass, never invents an input.

Reports which names are sitting on YOUR desk, batched by market, so a pack goes out once rather than a name at a time.

WHAT THE BOARD IS, AND WHY

engine/recalibration_run.py. Nothing is marked done and no step pointer is stored: each step is a PROBE over the artefacts that step produces, so a name sits wherever the probes stop.

A board kept by hand goes stale the first busy afternoon — silently.

ALL NINETY NAMES HAVE A DELIVERED STUDY. Twenty-three also commit a record the gates can read. A missing record is not a missing study, and reading it as one is what put "67 built from nothing" into the first version of this plan.

WAVE 1: the 23 record-backed re-issues, EGX first. Then Phase 1. Then the 67.

WHY THE BOARD EXISTS, IN ONE SENTENCE: a probe that comes back empty has not found nothing — the first hypothesis is that the probe did not run. That is how sixty-seven delivered studies came to be reported as absent, and it is why every step is now read off an artefact rather than off anyone's memory.