

Independent Valuation Study — Educational Analysis Saudi Telecom Company (Tadawul: 7010)

Fundamental analysis · Technical analysis · Monte Carlo simulation — one integrated read

Anchor: SAR 43.86 (2026-09-03 close) · 4,993.0 mn shares (issued capital over par, less treasury) · market capitalisation SAR 219.0 bn · the Kingdom's incumbent operator and the largest MENA telecom by market value: **stc KSA** (consumer, enterprise, wholesale — ~57% mobile share, the national fibre and 5G backbone) plus **subsidiaries** (solutions by stc 79%, stc bank 85%, stc Kuwait 51.8%, stc Bahrain, center3 data centres, sirar, iot squared) and **minority stakes** — 43.06% of the PIF-controlled tower company (TAWAL/Digital Infrastructure) and 9.97% of Telefónica · reporting currency SAR (pegged to USD at 3.75) · prices and probabilities computed 7 Jul 2026 from the attached daily history · primary lens: going-concern FCFF DCF, cross-checked by the dividend-policy DDM, relative multiples and normalized earnings · the swing factors: capex intensity against the locked SAR 2.20 dividend, and the discount-rate (beta) question.

READ FIRST — what this document is, and is not.

This study is a valuation exercise and an expression of personal analytical opinion, published free of charge for educational purposes: it shows how one analyst applies fundamental, technical and probabilistic methods to a listed company, and invites scrutiny of that methodology. It is NOT investment advice, NOT a recommendation or solicitation to buy, sell or hold any security, and NOT directed at the circumstances of any reader. The preparer is not licensed by any securities regulator in any jurisdiction, holds no brokerage or advisory authorisation, provides no financial consultancy, manages no money, and accepts no fees, funds or clients. See the Disclosure & Disclaimer at the end.

All values are model outputs presented as ranges and distributions because no single number should be relied on. Reported financials are the company's own disclosure (FY2023-FY2025 IR releases on the restated continuing-operations basis; Q1-2026 release, 28 Apr 2026; Q1-2026 interim financial statements) — all from stc.com. Forward-looking inputs — the segment growth and margin paths, capex intensity, the cost of capital, terminal growth, the multiples and the Monte-Carlo factor probabilities — are the preparer's judgments and are flagged throughout. Some balance-sheet detail lines are grouped estimates tying to disclosed totals. Consult a licensed financial advisor before any investment decision.

Headline

The model's read: the central sits below the traded price, and the study is HELD rather than published while it does. At SAR 43.86 — the latest known close, on 2026-09-03 — the shares sit about 13% above the central of SAR 38.14, which is the CASH-FLOW lens and not an average of the reads beside it. Those cross-checks do NOT cluster: the dividend-policy read discounts to 35, the enterprise multiple on this company's own history gives 48, a normalised-earnings read marks 40, and disclosed book value is 17 — a published range of SAR 37.4 to 48.0, and the study publishes that disagreement rather than averaging it away. The terminal carries 71% of the cash-flow value — disclosed, not buried — and is built on the 20.9-year asset life this company's own accounting-policies note discloses, never on the reciprocal of an inflation rate. Over three months the price map places the 5th-to-95th percentile band at roughly SAR 39-52, the median at 45.3, and the middle half at SAR 43-48.

Valuation summary — every read at a glance

One table for the four reads that follow — what the business is worth (fundamental), what the tape is doing (technical), where price could travel over three months (Monte Carlo), and how three independent expert methods land. Every row is developed in the sections and appendices below.

Lens / read	What it measures	Output	Takeaway
FUNDAMENTAL — what the business is worth (the anchor)			
FCFF DCF (primary)	Core operations + stake marks – net debt – NCI	SAR 38	-13% vs spot
Dividend discount (policy lens)	The locked SAR 0.55/quarter, discounted at K_e	SAR 35	-19%
Relative (EV/EBITDA)	FY26E EBITDA × justified 9.0×, bridged to equity	SAR 48	+9%
Normalized earnings	Ex-one-off PAT × through-cycle P/E	SAR 40	-9% · the floor
THE CENTRAL — the cash-flow lens, not an average	the class primary; the others are cross-checks	SAR 38.14	-13.0% vs SAR 43.86
Published range	the span of the present-value reads	SAR 37.4-48.0	
TECHNICAL — what the tape is doing (timing, not value)			
Trend & momentum	Price vs the 20/50/100/200-day averages	Above all four, by 1.8% on the highest	Firm, low energy
Momentum / range	RSI · MACD · 52-week range	RSI 66 · MACD +0.09 · 40.2-45.4	Firm, not overbought
THE PRICE MAP — where price could go, from the cone anchor			
1 month	50,000 paths, the production engine	p5 42 · p50 44.9 · p95 48	Median +0.5% on the anchor
3 months	same engine, the longer horizon	p5 39 · p50 45.3 · p95 52	Median +1.3%; a small upward lean
EXPERT PANEL — three independent methods (Appendix C)			
Expert 1 — cash returns / economic profit	ROIC vs WACC with fading excess returns	SAR 37	Most conservative
Expert 2 — normalized earnings power	Mid-cycle EPS × multiple	SAR 40	The floor-setter
Expert 3 — macro-policy scenario tree	Rate path & payout scenarios into a DDM	SAR 38	Closest to spot
Panel range	Spread = the fade-vs-franchise question	SAR 37-40	Centres ~SAR 38

Bottom line. The fundamental reads span SAR 37.4 to 48.0 and the central — the cash-flow lens, which is what this class is valued on — is SAR 38.14, 13.0% below the latest known price. THAT GAP IS WHY THIS STUDY IS HELD: a central more than a tenth below the traded price is a high-prior-of-defect region, and an eight-heading review is what stands between the number and a reader. The review did not move the answer toward the price and was not meant to — what it found were defects in the model, each corrected on its own evidence, and the answer moved where the corrections took it. The cross-checks

are published beside the central and never averaged into it: two of them are not permitted cross-checks for this class at all, and in the delivered edition they carried 45% of a weighted answer.

Company overview — stc at a glance

Item	Value
Listed entity	Saudi Telecom Company (stc Group), Tadawul: 7010
What it is	The Kingdom's incumbent telecom operator and MENA's largest telecom by market value; consumer, enterprise and wholesale connectivity plus a digital-infrastructure and fintech portfolio
Spot / date	SAR 43.86 · 2026-09-03 close, the latest known
Shares · market capitalisation	4,993.0 mn (issued capital over par, less treasury) · SAR 219.0 bn
FY2025 revenue / EBITDA / net profit	SAR 77,819 mn (+2.5%) · SAR 24,469 mn (31.4% margin) · SAR 14,828 mn (+12.5% adjusted; reported –39.9% vs FY24's TAWAL-gain year)
Q1-2026 revenue / net profit	SAR 19,939 mn (+3.8%) · SAR 3,696 mn (+12.0% ex non-recurring; +1.3% reported)
Segment split (FY2025)	stc KSA SAR 51,119 mn (consumer 32,826 · enterprise 13,514 · wholesale 4,779) · subsidiaries net ~SAR 26,700 mn
Balance sheet	Net debt SAR 111 mn at FY25 (~0.0× EBITDA); SAR 7.1 bn at Q1-26 after the Jan-2026 \$2 bn sukuk (~0.3×); cash framings: FS 21.4 bn incl. stc bank / IR core 15.4 bn
Dividend	SAR 0.55/quarter locked through Q3-2027 (SAR 2.20/yr ≈ 5.0% yield); FY24 also paid a SAR 2.00 special (cash paid 2025: SAR 4.20/sh ≈ 9.6%)
Key stakes	43.06% Digital Infrastructure Co (TAWAL + Zain towers; PIF 54%) · 9.97% Telefónica (€2.1 bn cost; ≈€1.96 bn market) · solutions by stc 79% (mkt cap ~SAR 26 bn) · stc bank 85% (8 mn customers)
52-week range	SAR 40.20 (1 Mar 2026) – 45.38 (30 Oct 2025)
Ownership	PIF 62.0% (after the Nov-2024 SAR 3.86 bn accelerated bookbuild) · free float ~38%
Corporate events	\$2 bn dual-tranche sukuk (Jan-2026, 4.489%/5.083%) · 26 mn-share ESIP buyback approved 7 May 2026 · center3-HUMAIN AI-data-centre framework (Dec-2025, toward 1 GW) · SilkLink Syria fibre corridor (SAR 3 bn, Feb-2026) · 2Q26 results due ~late Jul 2026

Source: stc FY2023–FY2025 IR releases, Q1-2026 release and interim FS, FY2025 earnings presentation (all stc.com); Saudi Exchange disclosures; PIF press releases. Values rounded.

1 Fundamental valuation

We value stc as a going-concern operator. ONE lens is the answer — a free-cash-flow-to-firm DCF — and the others are published beside it as cross-checks rather than averaged into it. The primary lens is that DCF, because a mature, capital-intensive network operator is ultimately worth the cash its infrastructure produces after the capex that keeps it competitive; the tower and Telefónica stakes sit outside the operating engine and are marked separately on the bridge. A dividend-discount read is the natural cross-check for a company whose board has

locked a SAR 0.55-per-quarter distribution through 2027. A relative EV/EBITDA read and a normalized-earnings read complete the set. The weights and the football field are in §1.5; the segment engine in §1.6; the crux — dividend cover against the capex cycle, in real units — in §1.7; the cost-of-capital build (every input sourced, both ERP bases published) in §1.8; and the sensitivity grids in §1.9. One lens-selection note, considered and set aside: stc does own a captive-finance arm in stc bank, but at SAR 2.0 bn of revenue — 2.5% of the group — and an early-stage loan book it does not yet warrant the split-legs treatment a bank leg would get; it is carried inside the subsidiaries line and flagged as such, to be revisited once its regulated financials give it a book worth marking separately. Throughout, the cost of equity is published explicitly as the normalised risk-free rate plus beta times the equity risk premium: $4.54\% + 0.7078 \times 5.72\% = 8.59\%$, on the market basis this study names as central, with the alternative basis published beside it in §1.8 rather than chosen silently. The beta behind it is a five-year weekly regression against the published index of the exchange the stock is listed on, not an assumed round number.

1.1 The FCFF DCF — the primary lens

The revenue engine is the §1.6 segment build: stc discloses revenue by segment rather than subscribers and revenue per subscriber, so the forecast grows the disclosed lines on their own drivers rather than manufacturing a bottom-up split it has no data for. Group revenue compounds 2.4% a year across the explicit window, from SAR 80,737 mn to 88,722 mn — consumer growing slowly, enterprise recovering as mega-project phasing normalises, wholesale on hosting and fixed-wireless backhaul, and the subsidiaries fastest of the four. The EBITDA margin runs 31.98% in the first forecast year and 31.82% in the last: essentially flat, and slightly DOWN rather than up. That is deliberate and it is the point of §1.6 — margins here are an OUTPUT of the cost build, not an input, so the model is not permitted to assume the mix improvement that a glide would represent. Capital intensity is 14.96% of revenue, held flat, and it is MEASURED rather than guided: it is the three filed years' own mean of 1.161 times the depreciation of the base being renewed. Management publishes a band and an earlier edition of this study took its path straight from it; a forward target leans the same way an optimistic model does, so guidance is scored against what happens and never consumed as an input. Working capital is projected from the asset-conversion cycle rather than plugged. Zakat and income tax enter at the 8.03% the three filed years imply on operating profit, with the disclosed reversal of a prior year's provision taken out rather than extrapolated.

SAR mn	FY26E	FY27E	FY28E	FY29E	FY30E
Group revenue	80,737	82,985	85,056	86,982	88,722
EBITDA	25,823	26,475	27,089	27,678	28,232
D&A	(10,407)	(10,697)	(10,964)	(11,212)	(11,437)
EBIT	15,416	15,778	16,125	16,466	16,795
NOPAT = EBIT × (1 – 8.03%)	14,177	14,510	14,829	15,143	15,446
+ D&A	10,407	10,697	10,964	11,212	11,437
– Capex	(12,081)	(12,417)	(12,727)	(13,015)	(13,276)
– Δ working capital	(-0)	(394)	(363)	(338)	(305)
FCFF	12,504	12,396	12,703	13,002	13,301
Discount factor	0.925	0.855	0.791	0.731	0.676
PV of FCFF	11,563	10,601	10,047	9,510	8,997

DCF bridge	SAR mn
Σ PV of explicit FCFF (FY26-30E)	50,719

Terminal value (perpetuity, g = 2.00%)	185,454
PV of terminal value	125,444
Enterprise value — core operations	176,163
Terminal value as % of enterprise value	71%
+ Associates & JVs (43.06% DIIC/TAWAL, carrying)	12,910
+ Telefónica 9.97% (market mark)	8,513
less: Net debt (IR basis, Q1-26) · NCI	(8,375) · (3,928)
Equity value	190,447
DCF fair value per share	SAR 38.14

Two honesty notes. First, 71% of the enterprise value sits in the terminal — at a 5.1-point WACC–g spread this is a duration bet dressed as a five-year model, which is why §1.9 sensitizes the WACC × g grid and the beta separately rather than hiding either. Second, the model FCFF (SAR 10.3 bn FY26E) runs ~SAR 2–3 bn richer than stc's own reported FY25 free cash flow (6.5 bn), because reported OCF absorbs receivables swings, early-retirement cash and zakat timing that a NOPAT-based FCFF smooths; Q1-26's FCF of 3.9 bn (+494% YoY) suggests the gap is closing, but Appendix A shows both series so the difference is visible rather than averaged away.

1.2 Dividend discount — the policy lens, as the cash-flow cross-check

stc is one of the few large emerging-market payers whose dividend is a stated, board-locked policy rather than a ratio: SAR 0.55 per quarter from the Q4-2024 distribution through the Q3-2027 distribution (announced 25 Aug 2024), with specials assessed quarterly on top — FY24 shareholders also received SAR 2.00. We discount the locked SAR 2.20 through 2027, step the DPS with earnings thereafter (payout ~75–80%), and grow the terminal dividend at 3.0%.

DDM build	FY26E	FY27E	FY28E	FY29E	FY30E
DPS (SAR)	2.20	2.20	2.30	2.40	2.55
PV of DPS @ Ke 7.90%	2.03	1.87	1.80	1.73	1.69

Σ PV of explicit dividends (FY26–30E)	SAR 9.11
Terminal DPS (FY31E) = FY30E × 1.03	SAR 2.63
Terminal value = TDPS / (Ke – g)	SAR 39.67
PV of terminal value	SAR 26.32
DDM fair value per share	SAR 35.43
Terminal value as % of the total	74%

DDM base ≈ SAR 35/share — a touch above spot: at a 7.9% cost of equity a locked 5.0% yield growing at ~3% is worth slightly more than the market pays, and any repeat of a special distribution (the quarterly-assessment clause) is free upside to this lens. The cross-check earns its place for a subtler reason: it is the one lens that cannot be flattered by the capex assumptions, because the board has pre-committed the cash.

1.3 Relative multiples

On trailing numbers stc trades at ~14.7× earnings, ~9.5× EV/EBITDA and 2.5× book — the premium name in a GCC cohort spanning 10.7× (Ooredoo) to 17.6× (du) on earnings. The

premium tracks the franchise: ~57% mobile share, the only national fixed/fibre footprint, net debt $\approx$ zero against regional peers at 1–2 $\times$ EBITDA, and the highest absolute yield commitment. The disciplined relative read is a justified EV/EBITDA on the FY26E build, bridged to equity with the same stake marks as the DCF.

Relative basis	Bear	Base	Bull
FY26E EBITDA (SAR mn)	25,746	25,746	25,746
Justified EV/EBITDA	8.0 $\times$	9.0 $\times$	10.0 $\times$
Fair value (SAR/share)	45.7	48.0	49.8
P/E cross-check at base (on FY26E EPS 2.82)		17.0 $\times$	

Relative base $\approx$ SAR 48 — in line with the DCF, which is itself informative: the market's multiple for this franchise already embeds roughly the same view of its cash engine as our explicit model.

1.4 Normalized earnings power — where this sits in the cycle

Cycle position first: unlike a developer or a smelter, stc's P&L has no violent cycle, but FY25 profit is still not a clean base — it carries a one-off SAR 466 mn zakat credit (prior-year provision reversals), and FY24 before it carried the SAR 12.9 bn TAWAL disposal gain, a SAR 1.5 bn withholding-tax reversal and a SAR 2.6 bn early-retirement charge. Margins sit mid-cycle: the forecast opens at an EBITDA margin of 31.98% against the filed FY2025 year, and it is an output of the cost build rather than a target, enterprise revenue is at the soft point of the government mega-project phasing, and the subsidiary portfolio (stc bank, center3) is still in its investment phase — 2–3 subsidiaries are guided to turn contribution-positive from 2026. Normalized attributable profit is therefore $\sim$ SAR 14.4 bn (reported FY25 14.8 bn less the zakat credit), or EPS $\approx$ 2.89, capitalized at a through-cycle 15 $\times$ — the stock's own multi-year median area and the peer-set mid.

Normalized-earnings basis	Bear	Base	Bull
Normalized PAT (SAR mn)	13,600	14,400	15,200
Justified P/E	13.5 $\times$	15.0 $\times$	16.5 $\times$
Fair value (SAR/share)	36.8	39.8	50.2

Normalized base $\approx$ SAR 40 — the floor of the set: it pays nothing for growth beyond today's earnings power and treats the AI-infrastructure build purely as cost.

1.5 Synthesis — the central, and the lenses beside it

We weight the DCF most heavily because it is the only lens that prices the full capex-and-recovery arc; the DDM carries the pre-committed cash; the relative lens anchors to what the market pays for GCC telecom cash flows today; normalized earnings is the ballast.

Lens	Weight	Bear	Base	Bull
FCFF DCF (primary)	35%	37	38.1	39
Dividend discount (policy)	25%	29	35.4	50
Relative	20%	46	48.0	50

(EV/EBITDA)				
Normalized earnings	20%	37	39.8	50
THE CENTRAL — the class primary		37.4	38.1	39.0




Figure 1 — Valuation football field. Bars span bear-bull per lens; the brass tick is each base case; the gold band is the range the present-value reads span; the ink line is the market price.

Central fair value ≈ SAR 38/share, -13% versus spot. The bear-bull span (37-48) is driven almost entirely by the DCF's terminal arithmetic — the bear case is WACC +100 bp with margins 50 bp lighter and capex 100 bp heavier; the bull is the mirror image. Note what the spread is NOT about: revenue. A ± 1 pp change to every growth rate moves the central by only ~SAR 2 — this is a margin, capex and discount-rate story, which is exactly what \$1.7 and \$1.9 sensitize.

1.6 The segments — a deeper look, and the driver table

Segment / leg	FY25 revenue	Trend	Margin role	Swing role
KSA Consumer (CBU)	SAR 32.8 bn (+3.4%)	Mobility +2.8%, fixed +6.6%; 30.6 mn subs (+5.3% Q1-26)	The cash cow	Competition watch
KSA Enterprise (EBU)	SAR 13.5 bn (+0.4%)	Government phasing; private sector +6%; flat ex-mega-projects	High-margin	Recovery lever
KSA Wholesale & Carrier	SAR 4.8 bn (+10.8%)	Hosting, FWA backhaul, national roaming (+32.6% national)	Structural	Steady
Subsidiaries (net)	~SAR 26.7 bn	solutions 12.7 bn · channels ~14.1 bn gross · stc bank 2.0 bn (+11%) · SCCC +62% · sirar +13% · center3	Dilutive today, guided to turn	The option book
Associates & stakes	off-P&L	43.06% towers (DIIC) · 9.97% Telefónica · iot squared 50%	Equity-method / marks	Bridge items

The build is per disclosed segment. THE PREVIOUS EDITION GREW FOUR BUSINESS UNITS — consumer, enterprise, wholesale and a subsidiaries residual — which is not how this company reports: note 9 discloses ELEVEN operating segments and each is grown on its own measured real rate. One of them, the Saudi operating business and two thirds of group revenue, is built as volume times price from the subscriber counts the earnings presentations disclose; the others are forecast on their net rate because no unit data is published for them, and that gap is stated rather than filled.

Operating segment (note 9)	FY24	FY25	FY26E	FY28E	FY30E
stc	49,644	51,119	52,708	55,066	57,340
Channels	15,111	14,085	14,254	14,534	15,058
Solutions	12,064	12,730	13,682	15,048	15,829
stc Kuwait	4,105	4,180	4,324	4,532	4,721
stc Bahrain	1,928	1,968	2,020	2,097	2,180
Center 3	1,912	1,962	2,057	2,192	2,291
stc Bank	1,262	1,401	1,590	1,875	2,001
Sirar	733	827	963	1,174	1,262
Specialized	372	355	345	334	343
iot2	301	322	337	359	376
SCCC	188	304	319	340	355
Other operating segments	729	65	68	73	76
Eliminations / adjustments	(12,455)	(11,499)	(11,930)	(12,568)	(13,110)
Group revenue	75,893	77,819	80,737	85,056	88,722
EBITDA margin			32.0%	31.8%	31.8%
Capital intensity (% of revenue)			15.0%	15.0%	15.0%

History: stc FY2025 earnings presentation and IR releases (stc.com), restated basis. Forecast drivers are the house's own flagged view. These are the company's own internal unit names, used here to describe what drives the business; the MODEL is built on the eleven operating segments note 9 discloses, which is a different and finer cut.

1.7 The crux — dividend cover against the capex cycle, in real units

Three judgments drive this valuation, and all three are observable rather than abstract. First and largest: capex intensity against the locked dividend. The policy dividend costs SAR 10.98 bn a year (2.20 x 4,993.0 mn shares); management guides capex to 15–17.5% of revenue with 2026–27 “edging up.” The cover table below is the whole tension in one place — at the top of the guided band the dividend is only 0.86x covered by model FY26E free cash flow and the balance sheet funds the rest; at the bottom it is fully covered. Each percentage point of capex intensity is worth ≈SAR 0.8 bn of annual FCF and ≈SAR 1.6 of DCF fair value per share.

Second: the discount rate. This is a 90%-equity-weighted WACC built on a 0.48 regressed beta from a nine-week window (\$1.8) — at $\beta = 1.0$ the DCF falls from ~50 to ~34, below spot, so \$1.9 publishes the full beta grid rather than letting one regression settle it. Third: the SAMA/Fed path — the rf is 5.5% today; each 50 bp off the Saudi curve adds ≈SAR 4–5 to the DCF and directly lowers the funding cost of the AI-infrastructure build. A fourth, slower variable: KSA mobile competition (Mobily's subscriber growth has outpaced stc's revenue growth for four quarters) — each 1 pp off consumer growth costs ≈SAR 2 of fair value.

FY26E scenario (real units)	Model FCF (SAR bn)	Dividend bill (SAR bn)	Cover
Capex at 13.6% (the	13.6	11.0	1.24x

lowest of the three filed years, 1.054x depreciation)			
Capex at 15.0% (adopted, the three-year mean of 1.161x depreciation)	12.5	11.0	1.14×
Capex at 16.1% (the highest of the three filed years, 1.252x depreciation)	11.6	11.0	1.05×

The dividend schedule and its stress test live in Appendix A.3. The dividend is policy-locked through the Q3-2027 distribution; the test is whether FCF or the balance sheet pays for it — at Q1-26 run-rate (FCF 3.9 bn vs a 2.74 bn quarterly dividend) it was FCF, for the first quarter in a year.

1.8 Macro and country — SAMA, oil, Vision 2030, and the sourced cost of capital

stc is a defensive claim on the Saudi macro, in three channels. Rates: SAMA shadows the Fed to defend the riyal peg (repo 4.25% / reverse repo 3.75% since 10 Dec 2025; the Fed held 3.50–3.75% on 17 Jun 2026), so the discount rate and the sukuk funding cost are set in Washington as much as Riyadh. Oil and the fiscal impulse: government ICT and giga-project spend (SAR ~32 bn of digital-government spend in 2025) drives the enterprise book — the FY25 softness was phasing, not demand. Vision 2030: the structural bid — data centres (the center3-HUMAIN 1 GW ambition inside a national AI push), 5G/fibre densification, and digital-services adjacencies — is what turns a utility growth profile into a utility-plus-options profile. Because the riyal is pegged, there is no currency-translation channel in the valuation and no currency factor in the price map. Every input in the cost-of-capital build is sourced and named:

Cost-of-capital build	Value	Source
Risk-free rate, as observed	5.52%	The sovereign yield on the house macro path for Saudi Arabia, carried with its own as-of date
— less this sovereign's own default spread	0.98%	Country risk is charged exactly once, and it is charged inside the equity risk premium below — so it is removed here rather than counted twice
Normalised risk-free rate	4.54%	The two lines above, by subtraction
Equity beta	0.7078	A 4.91-year weekly regression of stc against the published index of the exchange it is listed on (raw_indices/SA/TASI.csv, as of 2026-08-18): 252 observations, R ² 30.2%, standard error 0.1029. This is the first tier of the house preference order, not a stopgap — but it explains 30% of the variance and no more, which is why \$1.9 prices the answer at every beta up to 1.2
Equity risk premium (market basis, adopted)	5.72%	The premium published for Saudi Arabia specifically, on the basis this study names as central; the alternative basis is published beside it rather than chosen silently
Cost of equity	8.59%	The normalised risk-free rate plus beta times the premium
Cost of debt, before tax	4.86%	Built from the company's own facilities, and above this sovereign's own yield,

		as a same-currency corporate borrower must be
Cost of debt, after tax	3.89%	At the 8.03% effective rate the three filed years imply
Weights (equity / debt)	90.3% / 9.7%	Market-value equity, never book: the latest known close of SAR 43.86 on 2026-09-03 times 4,993,024 thousand shares outstanding — issued capital of SAR 50,000,000 thousand in shares of SAR 10 each, which divides to the 5,000,000 thousand shares note 17 of the reviewed 30 June 2026 interim itself states, less 6,976 thousand held in treasury. Debt is the borrowings total from the LATEST DISCLOSED balance sheet, 30 June 2026, SAR 23,536,554 thousand, which reproduces from the same interim's own financing cash flows to twelve basis points.
Weighted cost of capital	8.13%	The explicit-window rate the forecast is discounted at
Terminal cost of capital	8.13%	Flat against the explicit window, and that is not an oversight: the riyal is pegged, so this economy is already at its terminal cost of capital by construction of the peg and there is no normalisation to glide toward
Terminal growth (nominal)	2.00%	Terminal inflation of 2.0% plus a stated real growth of 0.0% — derived from the house macro path rather than chosen, and the real component is written down as the number it is

Two honesty notes on this build. First, the sovereign quote behind the risk-free rate is older than the fourteen-day bound this house sets: The Saudi sovereign quote on the house macro path carries an as-of date of 31 July 2026 and this schedule is struck on 5 September 2026, so it is 36 days old against the 14-day bound. It is accepted DELIBERATELY and the age is disclosed here rather than used quietly; refreshing the path is a macro-path task that moves every Saudi study at once, not a lever inside this one. Second, the beta is a proper five-year weekly regression against the exchange's own published index — the standard, not a stopgap — but its R^2 of 30% means roughly 70% of what moves this stock is specific to it rather than to the market. A beta is a statement about co-movement, and on this name that statement is weak. The grid in §1.9 prices the answer at every beta up to 1.2 for exactly that reason. Both are recorded with the study's own inputs.

1.9 Sensitivity — the margin, the capex, the rate spread, and the beta

The first grid re-prices the cash-flow model across the two real-unit operating levers (EBITDA margin and capital intensity); the second across the cost of capital $\times$ terminal growth; the third across the beta — which is a proper five-year weekly regression against the exchange's own index, but explains less than a third of this stock's variance, so it is the input whose uncertainty is widest even though its construction is sound.




Figure 2 — Fair value from the cash-flow model (SAR/share) across shifts in the EBITDA margin and in capital intensity. The bold cell is the base case. EVERY cell in the grid sits below the market price of SAR 43.86: on this model the disagreement with the market survives the whole plausible range of both operating levers, so it is not an artefact of where either one was set.

WACC \ terminal g	1.5%	2.0%	2.5%	3.0%	3.5%
7.13%	44.0	44.6	45.4	46.3	47.4
7.63%	40.7	41.1	41.5	41.9	42.5
8.13% (base)	38.0	38.1	38.3	38.5	38.7
8.63%	35.6	35.6	35.6	35.6	35.6
9.13%	33.6	33.5	33.4	33.2	33.0

Fair value from the cash-flow model (SAR/share) across the cost of capital and terminal growth. The base cell is 8.13% against 2.00%, giving SAR 38.14. The CDS-ERP alternative WACC (7.90%) sits between the third and fourth rows.

Beta	Cost of equity	Cost of capital	Value per share	Note
0.30	6.26%	6.03%	55.48	
0.50	7.40%	7.06%	45.18	
0.71	8.59%	8.13%	38.14	the adopted regression
0.85	9.40%	8.87%	34.59	
1.00	10.26%	9.64%	31.57	the house fallback, had the regression failed its usability gate
1.20	11.40%	10.67%	28.38	

The grid is mandatory disclosure here: the regression is a proper five-year weekly one against the exchange's own published index, but it explains 30% of this stock's variance and no more, so the answer is priced across the plausible range rather than at one point. The read survives moderate beta doubt — at 0.85 the value is SAR 34.59 — but not a full reversion to 1.0, where it falls to SAR 31.57, -28% against the market price.

2 Technical and price structure

The tape is firm and the stock is trading above its whole moving-average stack. At SAR 44.72 the price sits above all 4 of the 20-, 50-, 100- and 200-day averages, which are themselves compressed into a band of SAR 0.61 (43.30 to 43.91) — so the stack is flat while the price has stepped clear of it by 1.8% above the highest of the four. RSI(14) at 66 is firm but short of overbought, and the MACD histogram is positive, with the line above its signal. The 52-week range is narrow for an emerging-market name — 40.20 to 45.38, a span of 13% — and realised volatility over the trailing year (12.9%) is among the lowest on the exchange. Price is up 1.8% over the last twenty sessions and 1.6% over the last sixty. This is a quiet advance rather than a trending breakout: the direction is up, the energy behind it is low.

Indicator	Reading	What it says
Price on 2026-08-31	SAR 44.72	71th percentile of the 52-week range
SMA 20 / 50	SAR 43.91 / 43.60	Price above both, by 1.8% and 2.6%
SMA 100 / 200	SAR 43.51 / 43.30	Stack compressed into SAR 0.61 — a long flat base
RSI (14)	65.7	Firm but short of overbought
MACD (12,26,9)	+0.31 line / +0.21 signal / +0.09 hist	Positive, with the line above its signal
52-week range	SAR 40.20 – 45.38	A span of 13% — narrow
Change over 20 / 60 sessions	+1.8% / +1.6%	A quiet advance
Realised volatility (252 sessions)	12.9%	Very low; the forward read used in §3 is 16.5%




Figure 3 — Price against the moving-average stack, the last 260 sessions.

For the probabilistic work this matters in one way: a low-volatility tape produces a genuinely narrow three-month cone — the 5th-to-95th band in §3 runs -12.1% to +16.7% around the anchor — so even a modest fundamental surprise can carry the price to the edge of the distribution. The technical read neither confirms nor contradicts the fundamental one, and it is not asked to: it is the shortest of the three lenses and speaks only to the weeks ahead.

3 A probabilistic price map

The probability read below opens the section: computed from the same 50,000 paths as everything that follows, it is a summary of the distribution, not an input to it. Both horizons are CALENDAR — one month and three months from the anchor, graded on a stated date rather than after a count of trading sessions.

The probability read (3 months, graded 30 Nov 2026)	
Chance the price finishes above the anchor	56%
Chance of +10% against -10% — the odds	16% vs 8% · 2.0 : 1
Median level, and its move	SAR 45.30 (+1.3%)
The 50% band (25th-75th)	SAR 42.9 – 47.9 (-4.2% / +7.0% of the anchor)
Chance of touching ±10% at any point	27% / 16%

How the cone is built. We simulate 50,000 price paths from the anchor (seed 42) with the house engine: the width comes from a cascade of this stock's own recent, medium and longer-run variation, projected forward over the horizon (annualised 16.5% at three months, 15.3% at one), then scaled by the calibration this market's whole panel was fitted to (1.056). The shape of the tails is a Student-t drawn at the same fitted setting. The two numbers trade off against each other, so neither is quoted as precise on its own — the honest object is the cone they jointly produce, and the record of how often it has held follows below.

The centre of the cone is not a view. It is the cost of carry — what money earns while it waits — plus one small lean, described next. The paths deliberately do not know the fundamental fair value of \$1: the two lenses are kept apart so that when they agree it is information rather than an echo. \$1 says what the business is worth; \$3 says where the price could travel from where it is.

The direction call. This market runs a measured momentum lean, and every covered name states the direction its own reading points — including the names where the reading is too weak to move anything. stc reads +0.30, which is past the 0.25 threshold below which no tilt is applied at all, so the call is UP and it carries weight: it shifts the centre of the three-month cone by +0.24% and the one-month cone by +0.13%. Those are small numbers on purpose. The lean is capped at the strength the evidence actually measured; past that point a bigger number is a deliberately worse forecast rather than a braver one.

How often these bands have actually held

Over 58 resolved three-month forecasts, the price finished inside the 90% band 88% of the time, against the 90% that band aims at — and inside the 80% and 50% bands 78% and 47% of the time.

The band is 1.04 times the width of a simple no-forecast rule's — the price of covering this name honestly. A wider band is not automatically a worse one: covering a name honestly costs width, and the number is published beside the record rather than turned into a pass mark.

No flag is raised, and that is the ordinary case rather than an omission: on a two-sided test at the 5% level, 88% inside a 90% band over 58 forecasts is what a cone that works looks like ($p = 0.52$). Had the bands run persistently narrow — price escaping more often than a 90% band should allow — or persistently wide, this study would say so here in those words. It says nothing because nothing has been earned either way.

The record is long by this study's own standard: 58 resolved three-month forecasts is enough history to judge this name on its own rather than falling back on the market's, and every one of them is a forecast that was published before its outcome was known and graded on a date fixed in advance.

Percentile map (SAR/share)

Horizon	Graded on	p5	p25	p50	p75	p95
1 month	30 Sep 2026	41.7	43.6	44.9	46.3	48.5
3 months	30 Nov 2026	39.3	42.9	45.3	47.9	52.2

Lead with the 50% band, not the tails: three months out, half of all paths finish between roughly SAR 43 and 48 (-4.2% / +7.0%); at one month the band is about SAR 44-46. The 5th-to-95th cone is context rather than a forecast, and on this name it is 33% wide from end to end — the direct consequence of the calm regime the width estimator is reading.




Figure 4 — The forward price cone over 3 months, from the anchor of SAR 44.72. The median carries at the cost of money plus the small momentum lean. The brass dashed line is this study’s fundamental central of SAR 38.14 — below every part of the cone, including its 5th percentile, which never falls under SAR 39.31. The two lenses disagree, and the study shows the disagreement rather than reconciling it.




Figure 5 — Where the price lands after 1 month.




Figure 6 — Where the price lands after 3 months.

Level-touch ladder

The chance that price reaches a level at any point before the horizon, rather than finishing there (a running maximum for levels above, a running minimum for those below):

Level (SAR)	By 1 month	By 3 months	Note
50	2%	20%	Round number; +11.8% on the anchor
48	11%	41%	Above the 52-week high
46	48%	72%	Just past the Oct-25 high zone
44	60%	75%	The anchor itself
42	12%	36%	The March base
40	1%	14%	The 52-week-low shelf
38	0%	4%	Stress zone
36	0%	1%	Crisis zone

4 Comparison of the lenses

Three lenses, read side by side and never averaged. The cash-flow model is the answer this study publishes — SAR 38.14 against a market price of SAR 43.86, -13.0% — and the others sit beside it as cross-checks rather than as weights in a blend. They do not agree, and the disagreement is the most useful thing on this page: the earnings-multiple read on this company's own trading history lands at SAR 47.97, +26% above the cash-flow answer, while the disclosed book value is SAR 17.02 — a floor, not a valuation. The technical picture is a firm tape trading above its whole moving-average stack. The price map says the market is likely to keep the stock within a few riyals of where it is over the quarter (50% band 43–48).

Lens	Role	Reads	Value per share
Discounted cash flow	THE CENTRAL	Below the market	SAR 38.14 (-13%)
Enterprise multiple on own history	Cross-check	Above the market	SAR 47.97 (+9%)
Book value	Disclosed floor	Far below — as a floor should be	SAR 17.02
Dividend discount	Cross-check	Below the market	SAR 35.43 (-19%)

Normalised earnings power	Cross-check	Below the market	SAR 39.75 (-9%)
Technical read	Context	Above the whole moving-average stack	See §2
The price map (three months)	Context	Near-even, a small upward lean	Median 45.3; 5th-95th 39-52

The envelope this study publishes is the RANGE of the present-value reads, not an average of them: SAR 37.37 to 47.97. No weights are applied anywhere, because a number produced by averaging several methods is a new method with parameters nobody tested — it imports the weakness of the weakest lens at whatever weight somebody typed.

Where the lenses disagree, and what would settle it. The gap between the cash-flow answer and the multiple read is not a modelling artefact — it is one question asked two ways. The multiple lens capitalises what this company earns today at what the market has paid for those earnings over the last three financial year ends; the cash-flow lens charges for the capital the company must keep spending to go on earning them. The whole difference sits in the capital expenditure cycle, and \$1.7 prices it in real units: at the lowest capital intensity of the three filed years the cash-flow answer is SAR 37.37, at the highest it is lower still. So the honest statement is that this study reads stc below its market price BECAUSE it charges the company for the investment phase it is currently in, and a reader who believes that phase ends sooner, or earns more than the model credits, should hold the multiple read instead — the study publishes both rather than choosing for them. What would settle it is disclosed economics on the data-centre build and two more years of capital intensity: neither is available today, and neither is invented here.

5 Catalysts to watch

- **2Q26 results (~late July 2026).** The first full quarter to show whether Q1's cash-flow inflection (FCF 3.9 bn, 1.4× the quarterly dividend) is a run-rate or a blip; watch enterprise revenue (mega-project phasing) and capex phasing.
- **Quarterly dividend declarations.** SAR 0.55 is policy through Q3-2027 — the information is in any SPECIAL distribution (the policy explicitly allows quarterly assessment; FY24 paid SAR 2.00). A special would mechanically re-rate the DDM lens; its absence through 2026 keeps the balance sheet funding the build.
- **SAMA/Fed decisions (Sep/Oct 2026).** Each 25-50 bp off the Saudi curve lowers this study's discount rate and the marginal cost of the sukuk-funded infrastructure build — the single most mechanical catalyst.
- **center3-HUMAIN milestones.** The 1 GW AI-data-centre ambition (framework Dec-2025) converts to value only as contracted megawatts and disclosed economics; any capacity/offtake announcement is the option going in-the-money.
- **TAWAL/DIIC monetization.** stc still owns 43.06% of the region's largest tower company; an IPO or stake sale would crystallize a mark this study carries at a conservative SAR 4.6 bn book. Speculation only — nothing announced.
- **KSA mobile competitive prints.** Mobily's subscriber momentum vs stc's +5.3%; a price response from any of the three would show first in consumer ARPU-proxy (CBU revenue per sub).
- **Telefónica.** The 9.97% stake (€2.1 bn cost) trades ≈7% below cost with a halved dividend; further deterioration is a mark-to-market drag, any Spanish-market recovery a small tailwind — the stake is ~3.5% of the DCF equity bridge.
- **Index / flow events.** PIF sold 2% in Nov-2024 at SAR 38.6 (float now ~38%); any further sell-down is a supply event, any float-driven index-weight uplift a demand event.

6 Reading the probability zones

Translating the three-month distribution into plain zones, anchored on the cone's own origin of SAR 44.72 on 2026-08-31 — which is the last session in the price library and NOT the 2026-09-03 close of SAR 43.86 the valuation is measured against. Two clocks, and the study says which is which:

Zone at 3 months	Range	Probability	What it would mean
Deep downside	< SAR 38	2%	A regional shock or a price war; below the 52-week low
Lower band	SAR 38-42	16%	Capital spending runs hot and dividend cover stays below 1x; the March base is retested
Around the anchor	SAR 42-46	39%	The status quo: the yield carries the return
Upper band	SAR 46-50	31%	Cover proven or rate cuts delivered; the 52-week high is pressed
Strong upside	> SAR 50	12%	A special dividend, or the data-centre build credited with something

The distribution leans mildly to the upside — the ladder above gives a 27% chance of tagging +10% at some point against 16% for −10% — but the mass is concentrated: 86% of all paths finish between SAR 38 and 50, and half of them inside the much narrower SAR 42.9-47.9. That concentration is what a calm tape produces mechanically; it is not a view about the company, and a shift in the volatility regime would widen the cone at the next re-strike without anybody deciding anything.

7 Caveats and what would change our mind

- **The dividend-vs-capex tension is the model.** The FY26E cover of 0.86-1.04× is an estimate built on management's own capex band and our margin glide; two consecutive quarters of FCF below the SAR 2.74 bn quarterly dividend bill would push the DDM lens toward its bear case and the balance sheet toward releveraging — watch the quarterly FCF line, not the payout announcements.
- **Terminal-value dependency.** 71% of the cash-flow model's enterprise value sits beyond the fifth forecast year, at a spread of 6.1 points between the terminal discount rate and the 2.0% terminal growth. §1.9 is the honest statement that this valuation is substantially a judgement about the long-run price of money in Saudi Arabia. The terminal charge for keeping the asset base whole rests on the 21-year average life the company's own accounts imply for its plant, not on a life chosen here — but a network is a bundle of assets with very different lives, and a single average is the coarsest honest way to state it.
- **The sovereign quote is older than the house bound.** The Saudi sovereign quote on the house macro path carries an as-of date of 31 July 2026 and this schedule is struck on 5 September 2026, so it is 36 days old against the 14-day bound. It is accepted DELIBERATELY and the age is disclosed here rather than used quietly; refreshing the path is a macro-path task that moves every Saudi study at once, not a lever inside this one.
- **The beta is a real regression, and it explains less than a third of the variance.** The 0.71 beta is a 4.9-year weekly regression of this stock against the published index of the exchange it is listed on (252 observations, standard error 0.10), which is the standard this house requires and not a stopgap. What it does NOT do is explain the stock: the regression accounts for 30% of the variance, so roughly 70% of what moves stc is specific to stc. A

beta is a statement about how a stock moves with its market, and on this name that statement is weak. §1.9 shows the valuation at every beta up to 1.2 for exactly that reason.

- **Competition is the slow leak.** stc's consumer growth (+3%) already trails Mobily's subscriber momentum; a price war in KSA mobile — three players, one regulator, heavy 5G capacity — is the scenario in which the margin glide reverses and every lens compresses toward the normalized floor.
- **The subsidiary portfolio is an investment phase, not yet a return.** stc bank (8 mn customers, SAMA-licensed Jan-2025), center3, SCCC and iot squared are guided to turn contribution-positive from 2026; if they stay dilutive through FY27, the 31.82% terminal-year margin this model carries is wrong by construction — and it is worth saying that this forecast assumes NO mix improvement at all, so the risk here is a margin that falls rather than one that fails to rise.
- **The stake marks are point-in-time.** Telefónica (9.97%) is marked at market (≈SAR 8.6 bn, ~7% below cost) and the tower stake at carrying value (SAR 4.6 bn) — the former moves daily, the latter is conservative against any transaction mark; together they are ~5% of the bridge.
- **OCF conversion.** Reported FY25 FCF (6.5 bn) ran well below model FCFF on receivables build and one-off cash items; we model an explicit conversion drag, but a government-receivables cycle that widens rather than narrows would make even the 0.86× bear cover optimistic.
- **Technical reminder.** A compressed, low-volatility tape can break either way; the narrow cone in §3 is a regime read, not a promise. The technical section is context, not a trigger.

Appendix A Financial statements

Consolidated figures as disclosed by stc — FY2023-FY2025 IR releases on the restated continuing-operations basis (TAWAL and Digital Infrastructure Co reclassified to discontinued operations), Q1-2026 release and interim FS — all from stc.com, per the study's sourcing rule. SAR million. The five-year forecast is the model build (companion Excel, formula-linked to Assumptions).

A.1 Income statement — 3-year historical + 5-year forecast (consolidated, SAR mn)

Line	FY23	FY24	FY25	FY26E	FY27E	FY28E	FY29E	FY30E
Revenue	71,777	75,893	77,819	80,737	82,985	85,056	86,982	88,722
Gross profit	34,740	37,326	37,700	39,678	40,716	41,685	42,605	43,457
EBITDA	22,445	23,951	24,469	25,823	26,475	27,089	27,678	28,232
Depreciation, amortisation and impairment	(9,284)	(9,525)	(10,031)	(10,407)	(10,697)	(10,964)	(11,212)	(11,437)
Operating profit (EBIT)	13,161	14,426	14,438	15,416	15,778	16,125	16,466	16,795
Cost of the early retirement programme	(863)	(2,577)	(824)	(1,421)	(1,451)	(1,480)	(1,510)	(1,540)
Finance income	1,482	1,718	1,276	952	972	991	1,011	1,031

Finance cost	(1,068)	(1,234)	(1,125)	(1,444)	(1,474)	(1,503)	(1,533)	(1,564)
Other income, associates and other gains	1,275	-198	957	not forecast	not forecast	not forecast	not forecast	not forecast
Profit before zakat and income tax	13,987	12,134	14,723	13,503	13,825	14,132	14,433	14,722
Zakat and income tax	(1,327)	(1,192)	466	(1,116)	(1,143)	(1,168)	(1,193)	(1,217)
Profit from continuing operations	12,660	10,943	15,189	12,387	12,682	12,964	13,240	13,505

The three filed years are note 9's own reconciliation of segment revenue to profit from continuing operations, and they foot to the riyal in every column. THREE LINES ARE NOT FORECAST AND THE TABLE SAYS SO RATHER THAN LEAVING A READER TO NOTICE: net other income, the share of associates and net other gains have no disclosed driver between them and were worth 957 million in FY2025 alone, so the projected profit is BELOW what the same company would report if they recurred. Zakat is the rate the three years imply on profit before zakat, 8.27 per cent, with the reversal of prior years' provision that note 33(a) names put back — carrying that reversal forward would read 5.02 per cent and assume the company keeps finding it has over-provided.

A.2 Balance sheet — the filed years, and why there is no forecast column

Line	FY23	FY24	FY25
Total assets	159,646	160,638	157,477
Cash and equivalents	28,138	30,755	15,080
Borrowings	(21,958)	(15,132)	(15,191)
Equity attributable to the parent	78,985	89,417	83,414
Non-controlling interests	2,530	3,069	3,482

THERE IS NO FORECAST BALANCE SHEET IN THIS EDITION AND THE REASON IS A DEFECT IN THE FILING RATHER THAN A GAP IN THE MODEL. The latest reviewed balance sheet does not add up in its own current column — four subtotals and both totals — while every prior-year column foots exactly. That was not assumed to be an extraction problem: the file was re-fetched from the company's own site, re-extracted, and re-read by optical character recognition off the page rendered at 300 dots per inch, and both routes return the same figures. A statement is used only if it foots against its own arithmetic, and solving for the six figures that would make it foot would be inventing them. The lines this study's valuation does use are corroborated instead by the same interim's CASH FLOW statement, to the riyal. The previous edition printed a forecast balance sheet whose gaps were filled with grouped estimates chosen to make a balance-check row read zero; a check that cannot fail is not a check, and it is not reprinted here.

A.3 Cash flow, the two free-cash-flow framings, and the dividend schedule

Cash flow (SAR mn)	FY23	FY24	FY25	FY26E	FY28E	FY30E
Operating cash flow (disclosed / model)	22,418	19,885	18,283	23,647	26,155	28,441

Capex (disclosed / model)	(9,790)	(11,927)	(11,795)	(13,359)	(14,000)	(14,006)
Free cash flow	12,628	7,959	6,488	10,287	12,156	14,435
Dividends paid (attributable)	(8,000)	(18,712)	(10,978)	(10,978)	(11,477)	(12,724)

FY23–FY25 disclosed (stc IR); FY24 dividends-paid include the FY24 quarterly schedule; calendar-2025 cash dividends ≈ SAR 20.9 bn including the SAR 2.00/share special. Model OCF is a NOPAT-based construct and runs richer than the disclosed series — the conversion gap (receivables, ERP cash, zakat timing) is modelled as an explicit drag and discussed in §1.1; both framings shown per house rule.

Dividend schedule	FY25A	FY26E	FY27E	FY28E	FY29E	FY30E
DPS declared (SAR)	2.20	2.20	2.20	2.30	2.40	2.55
Dividend bill (SAR bn)	11.0	11.0	11.0	11.5	12.0	12.7
Payout of attributable NP	74%	78%	75%	75%	74%	74%
Yield at spot (declared)	5.0%	5.0%	5.0%	5.3%	5.5%	5.9%
Stress check	At the 17.5% top of the capex band, FY26E cover is 0.86× — the SAR 15.4 bn core cash pile funds the gap for >3 years before leverage exceeds 1× EBITDA; the policy is safe to Q3-2027 barring a margin break.					

Two yield framings (house rule): declared-regular 2.20/sh = 5.0%; cash paid during calendar-2025 incl. the FY24 special = 4.20/sh = 9.6%. A forecast-versus-actual scorecard and a new-versus-old reconciliation: not applicable on an initiation; both become standing sections from the first 3-month update.

Appendix B Peer frame, risk register, and the research register

B.1 The peer frame and the sector

stc is the premium name in a three-player Saudi market inside a GCC cohort where telecom equity is increasingly a yield-plus-digital-infrastructure asset class.

Name	P/E (t)	Div yield	One-line profile
-------------	----------------	------------------	-------------------------

stc (7010)	~14.7×	5.0% (9.6% incl. special)	Incumbent; ~57% mobile share; net debt ~0; AI-DC and fintech options
Mobily (7020)	~13.5×	4.5%	No.2; e& anchor (~28%); fastest subscriber momentum
Zain KSA (7030)	~12.8×	4.9%	No.3; tower-light; balance-sheet repair done
e& (UAE)	~13.2×	5.1%	UAE incumbent + Vodafone/PPF international portfolio
Ooredoo (Qatar)	~10.7×	5.7%	Multi-market; MENA data-centre pivot
du (UAE)	~17.6×	5.5%	No.2 UAE; hyperscale data-centre momentum — the multiple stc's DC build aspires to
Omantel / Beyon / Telecom Egypt	11.5× / 11.1× / 8.8×	3.9% / 7.1% / 1.5%	Regional context

Multiples approximate, mixed as-of dates (May-Jul 2026), secondary aggregators — context only, never model inputs.

Sector structure. Saudi mobile is a disciplined three-player market (~57/27/16 stc/Mobily/Zain) under CST regulation, with spectrum freshly allocated (Nov-2024 auction: stc took 600 MHz + 3.8 GHz), 5G at 63% populated coverage for stc and FWA adoption among the highest globally — FWA is both an opportunity (4.1 mn of stc's 6.0 mn fixed lines) and the competitive vector through which mobile capacity attacks fixed pricing. Fibre stays an stc moat (3.75 mn FTTH, 258k km). The Kingdom's AI push (HUMAIN, sovereign compute, the center3 1 GW ambition) is turning telecom capex into a national-strategy line item. Principal risks: a mobile price war; capex overshoot on the AI build; government-receivables cycles; subsidiary execution (stc bank credit costs as it scales); the Telefónica mark; regional geopolitics; and the rate path staying higher for longer.

B.2 Risk register

Risk	How it would show up first	What it is worth
A price war in Saudi mobile	Consumer revenue per subscriber, before it reaches the margin	The single largest downside: the margin glide is the model
Capital expenditure overshoot on the data-centre build	Capital intensity against the three filed years (1.05x to 1.25x depreciation)	The bear-to-bull span of the central: SAR 37.37 to the cash-flow bull
Government receivables cycles	Operating cash conversion, not the income statement	A cash-flow timing risk rather than a value risk, unless it persists
Subsidiary execution as the portfolio scales	Contribution turning positive on the guided schedule, or not	The margin glide to the terminal depends on it
The rate path staying higher for longer	The discount rate directly	71% of enterprise value sits beyond year five
The overseas listed stake	A mark that moves daily and is carried at market	A few per cent of the bridge; disclosed, not modelled

B.3 The research register — what was searched, and what came back empty

Before any forecast driver was set, the research ran in four rings — the world, the country, the industry, and the company — and every finding was recorded with its source and that source's own date. The

table below is that register in summary; the standalone bibliography carries every line of it, together with the full register of the 176 inputs this study consumes, each with its value, its source and its date.

Ring	Findings	Of which searches that came back empty	The sources they rest on
Global	4	1	company ir, company official, regulator official, search
Country	3	1	company official, regulator official, search
Industry	5	2	company ir, search
Company	7	0	company ir, company official
All four	19	4	sweep dated 2026-09-05

A NEGATIVE RESULT IS A RESULT, and this study records four of them. They are printed here rather than left out, because a question asked and answered with nothing is evidence, while a question never asked looks identical to one that found nothing:

- **Global ring — trade / sanctions / supply chains.** Searched the FY2025 audited statements and the FY2025 and H1-2026 earnings presentations for sanctions, export controls, supply-chain disruption or tariffs. Two hits in the statements, neither a risk disclosure: one lists 'supply chain and other related services' among a subsidiary's permitted activities, and one records a 2017 Kuwaiti court ruling on a regulatory tariff decree. No supply-chain or sanctions exposure is disclosed anywhere in the three documents.
- **Country ring — fiscal / political events with sector read-through.** Searched the FY2025 audited statements and both 2026 interim sets for any fiscal or political event with a stated effect on this company — a change in the zakat or income-tax basis, a subsidy, a state directive. The tax note records ordinary zakat and income tax and a FY2025 CREDIT of SAR 466,436 thousand against a FY2024 charge of 1,191,564, a swing of 1.66bn that the statements do not attribute to a policy change. Nothing else is disclosed.
- **Industry ring — new entrants (named-competitor level).** Searched the FY2025 audited statements and the FY2025 and H1-2026 earnings presentations for any named competitor, market-share figure or new entrant. ZERO hits across all three documents for 'competitor', 'market share', 'Mobily' or 'Zain'. This company does not disclose its competitive position in any document this study holds, and no peer multiple or share figure is therefore used anywhere in the model.
- **Industry ring — competitor capacity / price moves (named).** The same search as the new-entrants category and with the same result: no named competitor appears in any of the three documents, so no competitor capacity or price move can be cited. Recorded as an evidenced absence rather than closed by moving a company-ring finding into this row.

The company's own investor-relations channel was attempted first, before any aggregator, and the attempt is logged whether or not it succeeded. It was reached on 2026-09-05 — and the route matters: four direct address guesses each returned the site's own error page under a success code, which reads as a working page to anything that only checks the status. The sitemap found it. Every historical figure in this study comes from the company's own audited statements or its own investor material, and none from a data vendor, a broker or a press report.

Appendix C The expert valuation panel

Every Testahil study closes with a panel of standing expert personas — drawn from the house Expert Persona Library, not invented for the occasion, so each accumulates a track record across studies and a quarterly update is a re-run, not a re-training. For stc we cast the telecom

trio from the library's coverage map — cash-returns (DCF + returns on capital against WACC), earnings-power (normalized through-cycle earnings + multiples), and macro-policy (scenario-weighted policy options) — labelled Expert 1 / 2 / 3. Each runs a genuinely different method, derives its fair value from shown workings, and states a falsification condition.

C.1 Expert 1 — cash returns: ROIC against the cost of capital

Worldview and tradition. A business is worth the cash it returns over its life, and it creates value only when each riyal of capital earns above its cost. He looks past accounting earnings to free cash flow and to the economic-profit spread — and he is temperamentally suspicious of capex programs described with the word “vision.”

When it works / fails. Best for capital-intensive businesses where returns on capital are the crux — precisely a telecom. Fails where reinvestment economics are genuinely improving (past ROIC misleads a new-moat story) — his risk here if the AI-data-centre build earns structurally above telecom returns.

Expert 1's economic-profit test	Value
Invested capital (parent equity 85.0 bn plus net debt 8.4 bn)	SAR 93.4 bn
FY26E operating profit after tax, over that capital	SAR 14.2 bn → 15.2%
Cost of capital (he accepts the \$1.8 build)	8.13%
Economic profit = (return – cost) × capital	SAR 6.6 bn/yr
Fade: excess returns decay 2.5%/yr toward the cost of capital	Multiple on economic profit 11.6x
Core enterprise value = capital + present value of the fading economic profit	SAR 170 bn
+ stakes less net debt less the minority, divided by 4,993.0 mn shares	→ SAR 36.9 per share

Sensitivity, and the swing really is the fade rate alone — the cost of capital has its own grid in \$1.9 and is held here: at 1% a year, a durable moat, his value rises to SAR 40.01; at 4%, where competition and technology churn eat the spread, it falls to SAR 34.65. Cross-examination: he tells Expert 2 that a 15× multiple on normalized earnings quietly capitalizes today's ROIC forever without charging for the capital that sustains it — his fade does explicitly what the multiple hides. He tells Expert 3 that scenario trees on the policy rate are fine, but the bigger lever is inside the company: each percentage point of capex intensity is SAR 0.8 bn of cash that either earns the spread or doesn't.

Verdict, falsification, market-implied. Fair SAR 36.9 (range 35–40) — the panel's conservative anchor, below spot: on his arithmetic the market already pays for the excess returns to persist a decade-plus. Falsified by disclosed data-centre economics showing contracted returns above telecom ROIC (the fade would then be too harsh), or by return on capital holding above 15% through FY2028 while capital intensity normalises. What the price implies: at SAR 43.86 the market discounts a gentler fade than his 2.5% a year, i.e. it believes in the moat slightly more than he does.

C.2 Expert 2 — normalized earnings power

Worldview and tradition. An operating company is worth a fair multiple of its sustainable, mid-cycle earnings power; peaks, troughs and one-off gains are noise to be stripped before anything is capitalized.

When it works / fails. Best for stable operating businesses with a track record — a fit for stc, whose underlying earnings have grown 12–13% (adjusted) for two straight years. Fails at structural breaks: if the subsidiary portfolio re-rates the growth profile, his through-cycle multiple is too low; if a price war breaks the margin, too high.

Expert 2's normalisation	SAR mn
FY2025 profit before zakat, as filed	14,723

less zakat at the rate the three filed years imply	(1,217)
less the minority interest	(273)
Normalised profit attributable	13,233
Divided by shares in issue (4,993.0 million)	SAR 2.65 per share
Justified through-cycle price-to-earnings	15.0x
Fair value	SAR 39.8 per share

Sensitivity (the swing is the multiple): every turn of the price-to-earnings ratio is worth exactly his normalised earnings per share, SAR 2.65 — so at 13.5x he reads SAR 35.8 and at 16.5x SAR 43.7. His published range of SAR 36.8 to 50.2 spans 13.9x to 19.0x. Cross-examination: he tells Expert 1 that a fade model is a multiple wearing a lab coat — the honest disagreement is the number. He tells Expert 3 that the dividend lens undervalues whatever the board chooses not to distribute.

Verdict, falsification, market-implied. Fair SAR 39.8 (range 37–50), -9% against the market. On clean current earnings the stock is priced above his read: the market pays 16.5x his normalised earnings against the 15.0x he will justify, and the difference is growth not yet printed. Falsified by two consecutive years of double-digit adjusted earnings growth, which would make his base stale, or by the margin glide reversing.

C.3 Expert 3 — macro-policy: the scenario tree

Worldview and tradition. In a policy-driven market, policy outranks fundamentals: the Fed/SAMA rate path, the oil-funded fiscal impulse, and sovereign strategic priorities (Vision 2030, the AI build, PIF's portfolio choices) move this stock more than management execution does. He prices the equity as a probability-weighted set of policy worlds, expressed through the dividend stream — the one cash flow a policy-anchored shareholder actually receives.

When it works / fails. Best where binary policy catalysts dominate — rate cuts, sovereign flows, national-champion capex mandates. Fails through false precision: the probabilities are judgments, and the tree can miss the branch that grows (his own flag: a KSA price war appears in nobody's policy scenario, yet would dominate all of them).

Scenario (through 2027)	Prob.	World	DDM value
Easing + special dividends	30%	Fed/SAMA cut 75–100 bp; cover proven; a special repeats	SAR 51
Base: policy held, dividend locked	45%	Gradual cuts; capex mid-band; SAR 2.20 through 2027 then +3%	SAR 35
Higher-for-longer + capex overrun	25%	No cuts to mid-2027; capex at 17.5%; no specials	SAR 28
Probability-weighted fair value			SAR 38.1

Sensitivity (swing = the scenario weights): shifting 10 points from base to bear moves him ≈SAR 1.5; his answer is more stable than either colleague's because the dividend floor does most of the work in every branch. Cross-examination: he tells Expert 1 that a fade rate is unknowable to the decimal while the policy calendar is published — model what is scheduled. He tells Expert 2 that a through-cycle multiple assumes a cycle; Saudi rates are pegged to a foreign central bank, so the local 'cycle' is imported and can stay dislocated from local fundamentals for years.

Verdict, falsification, market-implied. Fair SAR 38.1 — the dividend model with its eyes open, and between his two colleagues, with Expert 2 the nearer to the market. Falsified by the board breaking the policy (either direction: a cut dividend or a step-change up), or by a Fed

path outside his tree (no cuts through 2027, or emergency easing). What the price implies: the market is pricing roughly his base case with a small weight on the bear — i.e. the locked dividend at a ~5% yield, and near-zero credit for specials or the AI build.

C.4 Cross-examination — each challenge conceded or rejected

Each expert puts a challenge to each colleague above. A challenge is worth nothing until somebody says whether it lands, so each is answered here, and two of the four are conceded.

Challenge	From → to	Answer
A fade model is a multiple wearing a lab coat — the honest disagreement is the number, not the machinery.	2 → 1	CONCEDED, partly. Both capitalise the same earnings; the fade makes the capital charge explicit and the multiple buries it. That is a real difference in what is visible, not in what is assumed — and it means Expert 1 owns his fade rate where Expert 2 does not own his multiple.
A through-cycle multiple assumes a cycle, and the Saudi rate cycle is imported from a foreign central bank.	3 → 2	CONCEDED. The riyal is pegged, so the local discount rate is set abroad and can stay dislocated from local fundamentals for years. Expert 2's multiple is a claim about the average of a cycle whose length he does not control.
Scenario trees on the policy rate are fine, but the bigger lever is inside the company: each point of capital intensity is real cash that either earns the spread or does not.	1 → 3	REJECTED, on this company's own numbers. The three filed years span a capital intensity of 1.05 to 1.25 times the depreciation of the base being renewed, and that whole span is worth less to the answer than the discount rate is. The policy lever is the larger one here; it is simply the one nobody controls.
The dividend lens undervalues whatever the board chooses not to distribute.	2 → 3	REJECTED, with a caveat. Retained cash shows up as future dividends in the same model, so it is not lost — but only if it earns its cost of capital on the way. That is Expert 1's question, not Expert 2's, and it is the one the data-centre build actually turns on.

Two conceded, two rejected. Nothing here changes any expert's number: a concession about what a method makes visible is not a concession about what it computes, and the panel is published with its disagreement intact rather than talked into a consensus.

C.5 The three in one room

Put the three in a room and the argument is about one thing: what happens to stc's return on capital as the Kingdom's AI-infrastructure build runs through its income statement.

Expert 1: "15 per cent returns on 93 billion of capital, fading as every telecom's returns have always faded. The data centres are capital expenditure with a press release until somebody shows me contracted economics. I pay SAR 37."

Expert 3: "Your fade rate is a guess dressed as physics. What is not a guess: the board has signed a cheque for SAR 2.20 a year through 2027, the central bank's next moves are the Federal Reserve's, and the sovereign has made this company its digital-infrastructure champion. Price the policy, not the physics — SAR 38."

Expert 2: "You are both reaching. Clean earnings are SAR 2.65 a share; the market pays 15 times that across the Gulf. Everything else — fades, scenario trees, gigawatts — is a story about the next turn of the multiple. SAR 40, and note that the market is paying 16.5 times, which is 10 per cent more than I will justify — so I am not agreeing with the price, I am naming the gap."

C.6 Reading the divergence




Figure C-1 — The three experts' fair-value ranges. Brass ticks are base cases; the gold band is the panel centre; the ink line is spot. The spread is the return-fade question.

Expert	Method	Single swing assumption	Base fair value
Expert 1	Cash returns / economic profit	The fade rate on excess returns (2.5%/yr)	SAR 36.9
Expert 2	Normalized earnings power	The through-cycle multiple (15x)	SAR 39.8
Expert 3	Macro-policy scenario tree	The scenario weights on the rate/payout path	SAR 38.1

The spread — SAR 36.9 to 39.8, 8% of the low — is narrow by this series' standards, and it measures one thing cleanly: how much of stc's current 15% return on capital survives the next decade of competition, technology churn and nation-scale capital spending. On that question the three barely disagree, and all three land below the market price: Expert 1 charges for the erosion explicitly, Expert 2 freezes today's clean earnings at a multiple he will defend, and Expert 3 prices the policy floor — three different routes to the same neighbourhood. The study's own central of SAR 38.14 sits inside that band, SAR 0.10 below the panel's own average, which is worth saying plainly: the cash-flow model and three independently-constructed methods arrive within a riyal or two of each other, and all of them below the price. An investor's position on this stock reduces to a position on one axis: if the data-centre build earns telecom-plus returns, the cash-flow model understates it; if it earns telecom-minus, Expert 1 is right and today's price already flatters it.

Each expert's point fair value, and its range where the method produces one, is recorded with the price it was struck against: SAR 43.86 on 2026-09-03.

About this series

Testahil publishes independent, educational valuation studies. Each is an attempt to reason transparently about what a security is worth, with every assumption shown and a companion model so readers can disagree productively. The house style is distributions, not tips: we describe ranges and probabilities, not targets, and we do not tell anyone what to do. Studies are framed as educational analysis, the preparer is not licensed by any securities regulator, and holdings are disclosed.

Disclosure & Disclaimer

Not investment advice. This document is educational and informational only. It is not, and must not be relied upon as, investment, financial, legal, accounting or tax advice, nor an offer, solicitation or recommendation to buy, sell or hold any security. It contains no price target and no rating.

No licence; no advisory relationship. The preparer is not registered or licensed with any securities or financial regulator in any jurisdiction — including the Saudi Capital Market Authority (CMA) — holds no brokerage or investment-advisory authorisation, and is not acting as your adviser or fiduciary. Nothing here is personalised to your circumstances.

Holdings disclosure. The preparer may hold, and may in the future take or dispose of, a position in the security discussed in this report, and may transact at any time without notice. This is a potential conflict of interest you should weigh.

Sources & accuracy. Reported financial and operating figures are drawn from the company's public disclosure (stc.com IR releases and interim financial statements) and other public sources believed reliable but not independently verified; they may contain errors or be superseded. Forward-looking inputs — the segment growth and margin paths, capex intensity, the derived risk-free rate and regressed beta, terminal growth, the multiples, the stake marks and the Monte-Carlo factor probabilities — are the preparer's own judgments and are inherently uncertain. Some balance-sheet detail lines are grouped estimates tying to disclosed totals.

Forward-looking statements. Any statements about the future are estimates subject to risks and uncertainties; actual results may differ materially. The Monte Carlo models price, not value, and encodes subjective probabilities for events that have not occurred.

No reliance; your responsibility. Do your own research and consult a licensed professional before making any decision. You are solely responsible for your investment decisions and their outcomes. To the maximum extent permitted by law, the preparer accepts no liability for any loss arising from use of this document.

Currency & figures. Figures are in Saudi riyals (SAR), millions unless stated; bn denotes billion. The riyal is pegged to the US dollar at 3.75. Rounding may cause totals to differ slightly. Spot price and market data are as of 7 July 2026 and change continuously.